




SAS[®] 9.4 Data Set Options: Reference, Fourth Edition



The correct bibliographic citation for this manual is as follows: SAS Institute Inc. 2016. *SAS® 9.4 Data Set Options: Reference, Fourth Edition*. Cary, NC: SAS Institute Inc.

SAS® 9.4 Data Set Options: Reference, Fourth Edition

Copyright © 2016, SAS Institute Inc., Cary, NC, USA

All Rights Reserved. Produced in the United States of America.

For a hard copy book: No part of this publication may be reproduced, stored in a retrieval system, or transmitted, in any form or by any means, electronic, mechanical, photocopying, or otherwise, without the prior written permission of the publisher, SAS Institute Inc.

For a web download or e-book: Your use of this publication shall be governed by the terms established by the vendor at the time you acquire this publication.

The scanning, uploading, and distribution of this book via the Internet or any other means without the permission of the publisher is illegal and punishable by law. Please purchase only authorized electronic editions and do not participate in or encourage electronic piracy of copyrighted materials. Your support of others' rights is appreciated.

U.S. Government License Rights; Restricted Rights: The Software and its documentation is commercial computer software developed at private expense and is provided with RESTRICTED RIGHTS to the United States Government. Use, duplication, or disclosure of the Software by the United States Government is subject to the license terms of this Agreement pursuant to, as applicable, FAR 12.212, DFAR 227.7202-1(a), DFAR 227.7202-3(a), and DFAR 227.7202-4, and, to the extent required under U.S. federal law, the minimum restricted rights as set out in FAR 52.227-19 (DEC 2007). If FAR 52.227-19 is applicable, this provision serves as notice under clause (c) thereof and no other notice is required to be affixed to the Software or documentation. The Government's rights in Software and documentation shall be only those set forth in this Agreement.

SAS Institute Inc., SAS Campus Drive, Cary, NC 27513-2414

April 2026

SAS® and all other SAS Institute Inc. product or service names are registered trademarks or trademarks of SAS Institute Inc. in the USA and other countries. ® indicates USA registration.

Other brand and product names are trademarks of their respective companies.

9.4-P9:ledsoptsref

Contents

<i>Syntax Conventions for the SAS Language</i>	<i>v</i>
<i>What's New in SAS 9.4 Data Set Options</i>	<i>xi</i>
Chapter 1 / About Data Set Options	1
SAS Compute Server and CAS	1
Definition of Data Set Options	1
Syntax	2
Using Data Set Options	2
Data Set Options Documented in Other SAS Publications	4
Chapter 2 / Dictionary of Data Set Options	5
Data Set Options by Category	6
Dictionary	9

Syntax Conventions for the SAS Language

Overview of Syntax Conventions for the SAS Language

SAS uses standard conventions in the documentation of syntax for SAS language elements. These conventions enable you to easily identify the components of SAS syntax. The conventions can be divided into these parts:

- syntax components
- style conventions
- special characters
- references to SAS libraries and external files

Syntax Components

The components of the syntax for most language elements include a keyword and arguments. For some language elements, only a keyword is necessary. For other language elements, the keyword is followed by an equal sign (=). The syntax for arguments has multiple forms in order to demonstrate the syntax of multiple arguments, with and without punctuation.

keyword

specifies the name of the SAS language element that you use when you write your program. Keyword is a literal that is usually the first word in the syntax. In a CALL routine, the first two words are keywords.

In these examples of SAS syntax, the keywords are bold:

CHAR (*string, position*)

CALL RANBIN (*seed, n, p, x*);

ALTER (*alter-password*)

BEST *w*.

REMOVE <*data-set-name*>

In this example, the first two words of the CALL routine are the keywords:

CALL RANBIN(*seed, n, p, x*)

The syntax of some SAS statements consists of a single keyword without arguments:

DO;

... SAS code ...

END;

Some system options require that one of two keyword values be specified:

DUPLEX | NODUPLEX

Some procedure statements have multiple keywords throughout the statement syntax:

CREATE <UNIQUE> **INDEX** *index-name* **ON** *table-name* (*column-1* <
column-2, ...>)

argument

specifies a numeric or character constant, variable, or expression. Arguments follow the keyword or an equal sign after the keyword. The arguments are used by SAS to process the language element. Arguments can be required or optional. In the syntax, optional arguments are enclosed in angle brackets (< >).

In this example, *string* and *position* follow the keyword CHAR. These arguments are required arguments for the CHAR function:

CHAR (*string, position*)

Each argument has a value. In this example of SAS code, the argument *string* has a value of 'summer', and the argument *position* has a value of 4:

```
x=char('summer', 4);
```

In this example, *string* and *substring* are required arguments, whereas *modifiers* and *startpos* are optional.

FIND(*string, substring* <*modifiers*> <*startpos*>)

argument(s)

specifies that one argument is required and that multiple arguments are allowed. Separate arguments with a space. Punctuation, such as a comma (,) is not required between arguments.

The MISSING statement is an example of this form of multiple arguments:

MISSING *character(s)*;

<LITERAL_ARGUMENT> *argument-1* <<LITERAL_ARGUMENT> *argument-2 ...* >

specifies that one argument is required and that a literal argument can be associated with the argument. You can specify multiple literals and argument

pairs. No punctuation is required between the literal and argument pairs. The ellipsis (...) indicates that additional literals and arguments are allowed.

The BY statement is an example of this argument:

BY <DESCENDING> *variable-1* <<DESCENDING> *variable-2* ...>;

argument-1 <*options*> <*argument-2* <*options*> ...>

specifies that one argument is required and that one or more options can be associated with the argument. You can specify multiple arguments and associated options. No punctuation is required between the argument and the option. The ellipsis (...) indicates that additional arguments with an associated option are allowed.

The FORMAT procedure PICTURE statement is an example of this form of multiple arguments:

PICTURE name <(format-options)>
<value-range-set-1 <(picture-1-options)>
<value-range-set-2 <(picture-2-options)> ...>;

argument-1=value-1 <*argument-2*=value-2 ...>

specifies that the argument must be assigned a value and that you can specify multiple arguments. The ellipsis (...) indicates that additional arguments are allowed. No punctuation is required between arguments.

The LABEL statement is an example of this form of multiple arguments:

LABEL *variable-1*=label-1 <*variable-2*=label-2 ...>;

argument-1 <, *argument-2*, ...>

specifies that one argument is required and that you can specify multiple arguments that are separated by a comma or other punctuation. The ellipsis (...) indicates a continuation of the arguments, separated by a comma. Both forms are used in the SAS documentation.

Here are examples of this form of multiple arguments:

AUTHPROVIDERDOMAIN (provider-1:domain-1 <, provider-2:domain-2, ...>
INTO :macro-variable-specification-1 <, :macro-variable-specification-2, ...>

Note: In most cases, example code in SAS documentation is written in lowercase with a monospace font. You can use uppercase, lowercase, or mixed case in the code that you write.

Style Conventions

The style conventions that are used in documenting SAS syntax include uppercase bold, uppercase, and italic:

UPPERCASE BOLD

identifies SAS keywords such as the names of functions or statements. In this example, the keyword ERROR is written in uppercase bold:

ERROR <message>;

UPPERCASE

identifies arguments that are literals.

In this example of the CMPMODEL= system option, the literals include BOTH, CATALOG, and XML:

CMPMODEL=BOTH | CATALOG | XML |

italic

identifies arguments or values that you supply. Items in italic represent user-supplied values that are either one of the following:

- nonliteral arguments. In this example of the LINK statement, the argument *label* is a user-supplied value and therefore appears in italic:

LINK *label*;

- nonliteral values that are assigned to an argument.

In this example of the FORMAT statement, the argument DEFAULT is assigned the variable *default-format*:

FORMAT *variable(s)* <format> <DEFAULT = *default-format*>;

Special Characters

The syntax of SAS language elements can contain the following special characters:

=

an equal sign identifies a value for a literal in some language elements such as system options.

In this example of the MAPS system option, the equal sign sets the value of MAPS:

MAPS=*location-of-maps*

< >

angle brackets identify optional arguments. A required argument is not enclosed in angle brackets.

In this example of the CAT function, at least one item is required:

CAT (*item-1* <, *item-2*, ...>)

|

a vertical bar indicates that you can choose one value from a group of values. Values that are separated by the vertical bar are mutually exclusive.

In this example of the CMPMODEL= system option, you can choose only one of the arguments:

CMPMODEL=BOTH | CATALOG | XML

...

an ellipsis indicates that the argument can be repeated. If an argument and the ellipsis are enclosed in angle brackets, then the argument is optional. The repeated argument must contain punctuation if it appears before or after the argument.

In this example of the CAT function, multiple *item* arguments are allowed, and they must be separated by a comma:

CAT (*item-1* <, *item-2*, ...>)

'value' or "value"

indicates that an argument that is enclosed in single or double quotation marks must have a value that is also enclosed in single or double quotation marks.

In this example of the FOOTNOTE statement, the argument *text* is enclosed in quotation marks:

FOOTNOTE <*n*> <*ods-format-options* 'text' | "text">;

;

a semicolon indicates the end of a statement or CALL routine.

In this example, each statement ends with a semicolon:

```
data namegame;
  length color name $8;
  color = 'black';
  name = 'jack';
  game = trim(color) || name;
run;
```

References to SAS Libraries and External Files

Many SAS statements and other language elements refer to SAS libraries and external files. You can choose whether to make the reference through a logical name (a libref or fileref) or use the physical filename enclosed in quotation marks.

If you use a logical name, you typically have a choice of using a SAS statement (LIBNAME or FILENAME) or the operating environment's control language to make the reference. Several methods of referring to SAS libraries and external files are available, and some of these methods depend on your operating environment.

In the examples that use external files, SAS documentation uses the italicized phrase *file-specification*. In the examples that use SAS libraries, SAS documentation uses the italicized phrase *SAS-library* enclosed in quotation marks:

```
infile file-specification obs = 100;  
libname libref 'SAS-library';
```

What's New in SAS 9.4 Data Set Options

Overview

SAS Data Set Options: Reference includes one new data set option and two enhanced data set options.

- The new [ENCRYPTKEY= on page 27](#) data set option specifies a key value. This new option enables you to use AES (Advanced Encryption Standard) encryption.
- The enhanced [ENCRYPT= on page 22](#) data set option supports AES encryption, which enables you to use the AES algorithm for stronger encryption.
- The default for the [EXTENDOBSCOUNTER= on page 30](#) data set option is now YES, which creates a SAS data file with an enhanced file format. The enhanced file format extends the observation count beyond the 32-bit long limitation.
- In SAS 9.4M5, stronger AES encryption was added with the AES2 key generation algorithm. ENCRYPT=AES2 specifies using this stronger algorithm.
- In SAS 9.4M5, the documentation has been merged with SAS Viya documentation. Functionality that is available in SAS 9.4M5 but not in SAS Viya is annotated.
- In the SAS Viya 3.3 release, the MAXTABLEMEM= data set option was added to specify the maximum amount of memory in bytes that each thread should allocate for in-memory blocks before converting to a memory-mapped file.
- In the SAS Viya 3.4 release, the WHERE= data set option is supported in CAS on input data sets and is not supported in CAS on output data sets.
- In SAS 9.4M7, the SESSION argument was added to the OUTREP= data set option.

About Data Set Options

<i>SAS Compute Server and CAS</i>	1
<i>Definition of Data Set Options</i>	1
<i>Syntax</i>	2
<i>Using Data Set Options</i>	2
Using Data Set Options with Input or Output SAS Data Sets	2
How Data Set Options Interact with System Options	3
<i>Data Set Options Documented in Other SAS Publications</i>	4

SAS Compute Server and CAS

In SAS Viya, data set options run in DATA steps in either Cloud Analytic Services (CAS) or the Compute Server. CAS is the cloud-based run-time environment that enables multi-threaded, parallel DATA step execution.

The Compute Server enables clients to submit data set options in SAS programs and stored procedures for processing by using the SAS language.

The syntax for running data set options in a DATA step or procedure in CAS is nearly identical to the syntax for running on the Compute Server. That is, CAS supports some of the data set options and provides much of the same functionality. However, not all data set options are supported for processing in CAS.

Definition of Data Set Options

Data set options specify actions that apply only to the SAS data set with which they appear. Data set options enable you to perform these operations:

- rename variables

- select only the first or last n observations for processing
- drop variables from processing or from the output data set
- keep variables in the data set
- specify a password for a data set
- specify encryption for a data set

Syntax

Specify a data set option in parentheses after a SAS data set name. Specify several data set options by separate them with spaces.

(option-1=value-1< ...option-n=value-n>)

The following examples show data set options in SAS statements:

```

■ data scores(keep=team game1 game2 game3);
■ data mydata(index=(b k) label='label for my data set' drop=p read=secret);
■ data new(drop=i n index=(j combo=(x1 a1 a20 b1 b50 )));
■ data idxdup2(compress=yes index=(ok1 ok2 ssn/unique ok3));
■ proc print data=new(drop=year);
■ set old(rename=(date=Start_Date));

```

Using Data Set Options

Using Data Set Options with Input or Output SAS Data Sets

Most SAS data set options apply to either input or output SAS data sets in DATA steps or procedure (PROC) steps. If a data set option is associated with an input data set, the action applies to the data set that is being read. If the data set option appears in the DATA statement or after an output data set specification in a PROC step, SAS applies the action to the output data set. In the DATA step, data set options for output data sets must appear in the DATA statement, not in any OUTPUT statements that might be present.

Some data set options, such as COMPRESS=, are meaningful only when you create a SAS data set because they set attributes that exist for the duration of the data

set. To change or cancel most data set options, you must re-create the data set. You can change other options (such as PW= and LABEL=) with PROC DATASETS. For more information, see [“DATASETS Procedure” in Base SAS Procedures Guide](#).

When data set options appear in input and output data sets in the same DATA or PROC step, SAS first applies data set options to input data sets. Then SAS evaluates programming statements or applies data set options to output data sets. Likewise, data set options that are specified for the data set that is being created are applied after programming statements are processed. For example, when you are using the RENAME= data set option, the new names are not associated with the variables until the DATA step ends.

In some instances, data set options conflict when they are used in the same statement. For example, you cannot specify the DROP= and KEEP= data set options for the same variable in the same statement. Timing can also be an issue in some cases. For example, if you are using KEEP= and RENAME= on a data set that is specified in the SET statement, KEEP= needs to use the original variable names. SAS processes KEEP= before the data set is read. The new names that are specified in RENAME= apply to the programming statements that follow the SET statement.

How Data Set Options Interact with System Options

Many system options and data set options share the same name and have the same function. System options remain in effect for all DATA steps and PROC steps in a SAS job or session.

The data set option overrides the system option for the data set in the step in which it appears. In this example, the OBS= system option in the OPTIONS statement specifies that only the first 100 observations are processed from any data set within the SAS job. However, the OBS= data set option in the SET statement overrides the system option for data set TWO. OBS= specifies that only the first five observations are read from data set TWO. The PROC PRINT step prints the data set FINAL. This data set contains the first five observations from data set TWO, followed by the first 100 observations from data set THREE:

```
options obs=100;

data final;
    set two(obs=5) three;
run;

proc print data=final;
run;
```

Data Set Options Documented in Other SAS Publications

In addition to data set options documented in *SAS Data Set Options: Reference*, data set options are also documented in the following publications:

- *SAS Companion for Windows*
- *SAS Companion for UNIX Environments*
- *SAS Companion for z/OS*
- *SAS National Language Support (NLS): Reference Guide*
- *SAS Scalable Performance Data Engine: Reference*
- *SAS/ACCESS for Relational Databases: Reference*
- *Base SAS Procedures Guide*
- *SAS Cloud Analytic Services: User's Guide*

Dictionary of Data Set Options

<i>Data Set Options by Category</i>	6
<i>Dictionary</i>	9
ALTER= Data Set Option	9
BUFNO= Data Set Option	11
BUFSIZE= Data Set Option	13
CNTLLEV= Data Set Option	14
COMPRESS= Data Set Option	16
DLDMGACTION= Data Set Option	19
DROP= Data Set Option	20
ENCRYPT= Data Set Option	22
ENCRYPTKEY= Data Set Option	27
EXTENDOBSCOUNTER= Data Set Option	30
FILECLOSE= Data Set Option	32
FIRSTOBS= Data Set Option	33
GENMAX= Data Set Option	35
GENNUM= Data Set Option	36
IDXNAME= Data Set Option	38
IDXWHERE= Data Set Option	40
INDEX= Data Set Option	42
IN= Data Set Option	44
KEEP= Data Set Option	45
LABEL= Data Set Option	47
MAXTABLEMEM= Data Set Option	49
OBSBUF= Data Set Option	50
OBS= Data Set Option	52
OUTREP= Data Set Option	61
POINTOBS= Data Set Option	64
PW= Data Set Option	66
PWREQ= Data Set Option	67
READ= Data Set Option	68
RENAME= Data Set Option	70
REPEMPTY= Data Set Option	73
REPLACE= Data Set Option	76
REUSE= Data Set Option	77
ROLE= Data Set Option	78
SORTEDBY= Data Set Option	79

SPILL= Data Set Option	82
TOBSNO= Data Set Option	90
TYPE= Data Set Option	91
USEDIRECTIO= Data Set Option: Linux	92
WHERE= Data Set Option	93
WHEREUP= Data Set Option	96
WRITE= Data Set Option	98

Data Set Options by Category

The categories for SAS data set options correspond to the SAS data set option groups:

Table 2.1 *Data Set Options by Category*

Category	Description
CAS	options that run on the CAS server
Data Set Control	options that are associated with data sets
Observation Control	options that are associated with observations
User Control of SAS Index Usage	options that are associated with indexes
Variable Control	options that are associated with variables
Miscellaneous	option that is associated with tape position

Category	Language Elements	Description
CAS	COMPRESS= Data Set Option (p. 16)	Specifies how observations are compressed in a new output SAS data set.
	DROP= Data Set Option (p. 20)	For an input data set, excludes the specified variables from processing; for an output data set, excludes the specified variables from being written to the data set.
	IN= Data Set Option (p. 44)	Creates a Boolean variable that indicates whether the data set contributed data to the current observation.
	KEEP= Data Set Option (p. 45)	For an input data set, specifies the variables to process; for an output data set, specifies the variables to write to the data set.
	LABEL= Data Set Option (p. 47)	Specifies a label for a SAS data set.

Category	Language Elements	Description
	MAXTABLEMEM= Data Set Option (p. 49)	Specifies the maximum amount of memory in bytes that each thread should allocate for in-memory blocks before converting to a memory-mapped file.
	RENAME= Data Set Option (p. 70)	Changes the name of a variable.
	REPLACE= Data Set Option (p. 76)	Specifies whether a new SAS data set that contains data can overwrite an existing data set.
	WHERE= Data Set Option (p. 93)	Specifies specific conditions to use to select observations from a SAS data set.
Data Access	MAXTABLEMEM= Data Set Option (p. 49)	Specifies the maximum amount of memory in bytes that each thread should allocate for in-memory blocks before converting to a memory-mapped file.
Data Set Control	ALTER= Data Set Option (p. 9)	Assigns an ALTER= password to a SAS file that prevents users from replacing or deleting the file, and enables access to a read- and write-protected file.
	BUFNO= Data Set Option (p. 11)	Specifies the number of buffers to be allocated for processing a SAS data set.
	BUFSIZE= Data Set Option (p. 13)	Specifies the size of a permanent buffer page for an output SAS data set.
	CNTLLEV= Data Set Option (p. 14)	Specifies the level of shared access to a SAS data set.
	COMPRESS= Data Set Option (p. 16)	Specifies how observations are compressed in a new output SAS data set.
	DLDMGACTION= Data Set Option (p. 19)	Specifies the action to take when a SAS data set in a SAS library is detected as damaged.
	ENCRYPT= Data Set Option (p. 22)	Specifies whether to encrypt an output SAS data set.
	ENCRYPTKEY= Data Set Option (p. 27)	Specifies a key value for AES (Advanced Encryption Standard) encryption.
	EXTENDOBSCOUNTER= Data Set Option (p. 30)	Specifies whether to extend the maximum observation count in a new output SAS data file.
	GENMAX= Data Set Option (p. 35)	Requests generations for a new data set, modifies the number of generations for an existing data set, and specifies the maximum number of versions.
	GENNUM= Data Set Option (p. 36)	Specifies a particular generation of a SAS data set.
	INDEX= Data Set Option (p. 42)	Defines an index for a new output SAS data set.

Category	Language Elements	Description
	LABEL= Data Set Option (p. 47)	Specifies a label for a SAS data set.
	OBSBUF= Data Set Option (p. 50)	Determines the size of the view buffer for processing a DATA step view.
	OUTREP= Data Set Option (p. 61)	Specifies the data representation for the output SAS data set.
	PW= Data Set Option (p. 66)	Assigns a READ, WRITE, and ALTER password to a SAS file, and enables access to a password-protected SAS file.
	PWREQ= Data Set Option (p. 67)	Specifies whether to display a password dialog box.
	READ= Data Set Option (p. 68)	Assigns a READ= password to a SAS file that prevents users from reading the file, unless they enter the password.
	REEMPTY= Data Set Option (p. 73)	Specifies whether a new, empty data set can overwrite an existing SAS data set.
	REPLACE= Data Set Option (p. 76)	Specifies whether a new SAS data set that contains data can overwrite an existing data set.
	REUSE= Data Set Option (p. 77)	Specifies whether new observations can be written to available space in compressed SAS data sets.
	ROLE= Data Set Option (p. 78)	Identifies the fact table for a star schema join.
	SORTEDBY= Data Set Option (p. 79)	Specifies how a data set is currently sorted.
	SPILL= Data Set Option (p. 82)	Specifies whether to create a spill file for non-sequential processing of a DATA step view.
	TOBSNO= Data Set Option (p. 90)	Specifies the number of observations to send in a client/server transfer.
	TYPE= Data Set Option (p. 91)	Specifies the data set type for a specially structured SAS data set.
	USEDIRECTIO= Data Set Option: Linux (p. 92)	Turns on direct file I/O for a library that contains the file to which the ENABLEDIRECTIO option has been applied.
	WRITE= Data Set Option (p. 98)	Assigns a WRITE= password to a SAS file that prevents users from writing to a file, unless the users enter the password.
Miscellaneous	FILECLOSE= Data Set Option (p. 32)	Specifies how a tape is positioned when a SAS data set is closed.
Observation Control	FIRSTOBS= Data Set Option (p. 33)	Specifies the first observation that SAS processes in a SAS data set.

Category	Language Elements	Description
	IN= Data Set Option (p. 44)	Creates a Boolean variable that indicates whether the data set contributed data to the current observation.
	OBS= Data Set Option (p. 52)	Specifies the last observation that SAS processes in a data set.
	POINTOBS= Data Set Option (p. 64)	Specifies whether SAS creates compressed data sets whose observations can be randomly accessed or sequentially accessed.
	WHERE= Data Set Option (p. 93)	Specifies specific conditions to use to select observations from a SAS data set.
	WHEREUP= Data Set Option (p. 96)	Specifies whether to evaluate new observations and modified observations against a WHERE expression.
User Control of SAS Index Usage	IDXNAME= Data Set Option (p. 38)	Directs SAS to use a specific index to match the conditions of a WHERE expression.
	IDXWHERE= Data Set Option (p. 40)	Specifies whether SAS uses an index search or a sequential search to match the conditions of a WHERE expression.
Variable Control	DROP= Data Set Option (p. 20)	For an input data set, excludes the specified variables from processing; for an output data set, excludes the specified variables from being written to the data set.
	KEEP= Data Set Option (p. 45)	For an input data set, specifies the variables to process; for an output data set, specifies the variables to write to the data set.
	RENAME= Data Set Option (p. 70)	Changes the name of a variable.

Dictionary

ALTER= Data Set Option

Assigns an ALTER= password to a SAS file that prevents users from replacing or deleting the file, and enables access to a read- and write-protected file.

Valid in: DATA step and PROC steps

Category: Data Set Control

Restriction: Use in the Compute Server, but not in CAS.

- Note: Check your log after this operation to ensure password security. For more information, see [“Blotting Passwords and Encryption Key Values” in SAS Programmer’s Guide: Essentials](#).
- See: ALTER= Data Set Option under OpenVMS, UNIX, or z/OS in the documentation for your operating environment.

Syntax

ALTER=*alter-password*

Syntax Description

alter-password

must be a valid SAS name. For more information, see [“SAS Words” in SAS Programmer’s Guide: Essentials](#).

Details

Overview of the ALTER= Option

The ALTER= option applies to all types of SAS files except catalogs. You can use this option to assign a password to a SAS file or to access a read-protected, write-protected, or alter-protected SAS file.

When replacing a SAS data set that is protected with an ALTER password, the new data set inherits the ALTER password. To change the ALTER password for the new data set, use the MODIFY statement in the DATASETS procedure.

Note: A SAS password does not control access to a SAS file outside of SAS. Use the operating system-supplied utilities and file-system security controls in order to control access to SAS files outside of SAS.

Restriction for Metadata-Bound Libraries

You cannot use this option with data sets that are bound to metadata-bound libraries. Metadata-bound libraries are password-protected, and you must specify that password to access data sets that are bound to them. To change this password, use PROC AUTHLIB. You cannot use PROC DATASETS to add a password to a data set that is bound to a metadata-bound library. For more information, see [“AUTHLIB Procedure” in Base SAS Procedures Guide](#).

See Also

- [“Manipulating Passwords” in Base SAS Procedures Guide](#)

Data Set Options:

- “ENCRYPT= Data Set Option” on page 22
- “PW= Data Set Option” on page 66
- “READ= Data Set Option” on page 68
- “WRITE= Data Set Option” on page 98

BUFNO= Data Set Option

Specifies the number of buffers to be allocated for processing a SAS data set.

Valid in: DATA step and PROC steps

Category: Data Set Control

Restriction: Use in the Compute Server, but not in CAS.

See: BUFNO= Data Set Option in the documentation for your operating environment.

Syntax

BUFNO= *n* | *nK* | *nM* | *nG* | *hexX* | MIN | MAX

Optional Arguments

n* | *nK* | *nM* | *nG

specifies the number of buffers in multiples of 1 (bytes); 1,024 (kilobytes). For example, a value of 8 specifies 8 buffers, and a value of 1k specifies 1,024 buffers.

Note: You can also use the KB, MB, and GB syntax notations.

hexX

specifies the number of buffers as a hexadecimal value. You must specify the value beginning with a number (0–9), followed by an X. For example, the value 2dx sets the number of buffers to 45.

MIN

sets the minimum number of buffers to 0, which causes SAS to use the minimum optimal value for the operating environment. This is the default.

MAX

sets the number of buffers to the maximum possible number in your operating environment, up to the largest 4-byte, signed integer, which is $2^{31}-1$, or approximately 2 billion.

Details

The buffer number is not a permanent attribute of the data set; it is valid only for the current SAS session or job.

BUFNO= applies to SAS data sets that are opened for input, output, or update.

A larger number of buffers can speed execution time by limiting the number of input and output (I/O) operations that are required for a particular SAS data set. However, the improvement in execution time comes at the expense of increased memory consumption.

To reduce I/O operations on a small data set as well as speed execution time, allocate one buffer for each page of data to be processed. This technique is most effective if you read the same observations several times during processing.

To request that SAS allocate the number of buffers based on the number of data set pages and index file pages, use the SASFILE global statement.

Operating Environment Information: The default value for BUFNO= is determined by your operating environment and is set to optimize sequential access. To improve performance for direct (random) access, change the value for BUFNO=. For the default setting and possible settings for direct access, see the BUFNO= data set option in the SAS documentation for your operating environment.

Comparisons

- If the BUFNO= data set option is not specified, then the value of the BUFNO= system option is used. If both are specified in the same SAS session, the value that is specified for the BUFNO= data set option overrides the value that is specified for the BUFNO= system option.

See Also

Data Set Options:

- [“BUFSIZE= Data Set Option” on page 13](#)

System Options:

- [“BUFNO= System Option” in SAS System Options: Reference](#)

Statements:

- [“SASFILE Statement” in SAS Global Statements: Reference](#)

BUFSIZE= Data Set Option

Specifies the size of a permanent buffer page for an output SAS data set.

Valid in:	DATA step and PROC steps
Category:	Data Set Control
Restrictions:	Use with output data sets only. This data set option is not supported in a DATA step that runs in CAS.
See:	BUFSIZE= Data Set Option under UNIX, z/OS, or OpenVMS in the documentation for your operating environment.

Syntax

BUFSIZE=*n* | *nK* | *nM* | *nG* | *hexX* | **MAX**

Syntax Description

n* | *nK* | *nM* | *nG

specifies the page size in multiples of 1 (bytes); 1,024 (kilobytes); 1,048,576 (megabytes); or 1,073,741,824 (gigabytes). For example, a value of **8** specifies a page size of 8 bytes, and a value of **4k** specifies a page size of 4,096 bytes.

Note: You can also use the KB, MB, and GB syntax notations.

Note: If the system option and the data set option are not set, the default is 0. As a result, SAS uses the minimum optimal page size for the operating environment. The BUFSIZE= system option is used in either of the following scenarios:

- if the BUFSIZE= data set option is not set
- if the BUFSIZE= data set option is set to zero

Use BUFSIZE=0 to reset the buffer page size to the default value in your operating environment.

hexX

specifies the page size as a hexadecimal value. You must specify the value beginning with a number (0–9), followed by an X. For example, the value **2dx** sets the page size to 45 bytes.

MAX

sets the page size to the maximum possible number in your operating environment, up to the largest 4-byte, signed integer, which is $2^{31}-1$, or approximately 2 billion bytes.

Details

The page size is the amount of data that can be transferred for a single I/O operation to one buffer. The page size is a permanent attribute of the data set and is used when the data set is processed.

A larger page size can speed execution time by reducing the number of times SAS has to read from or write to the storage medium. However, the improvement in execution time comes at the expense of increased memory consumption.

To change the page size, use a DATA step to copy the data set and either specify a new page or use the SAS default. To reset the page size to the default value in your operating environment, use BUFSIZE=0.

Note: You can use the COPY procedure to copy a data set to another library that is allocated with a different engine. The specified page size of the data set is not retained.

Operating Environment Information: The default value for BUFSIZE= is determined by your operating environment and is set to optimize sequential access. To improve performance for direct (random) access, change the value for BUFSIZE=. For the default setting and possible settings for direct access, see the BUFSIZE= data set option in the SAS documentation for your operating environment.

See Also

Data Set Options:

- [“BUFNO= Data Set Option” on page 11](#)

System Options:

- [“BUFSIZE= System Option” in SAS System Options: Reference](#)

CNTLLEV= Data Set Option

Specifies the level of shared access to a SAS data set.

Valid in:	DATA step and PROC steps
Category:	Data Set Control
Restrictions:	Use in the Compute Server, but not in CAS. Specify for input data sets only.

Syntax

CNTLLEV=LIB | MEM | REC

Syntax Description

LIB

specifies that concurrent access is controlled at the library level. Library-level control restricts concurrent access to only one update process to the library.

MEM

specifies that concurrent access is controlled at the SAS data set (member) level. Member-level control restricts concurrent access to only one update or output process to the SAS data set. If the data set is open for an update or output process, then no other operation can access the data set. If the data set is open for an input process, then other concurrent input processes are allowed, but no update or output process is allowed.

REC

specifies that concurrent access is controlled at the observation (record) level. Record-level control allows more than one Update access to the same SAS data set, but it denies concurrent update of the same observation.

Details

The CNTLLEV= option specifies the level at which shared Update access to a SAS data set is denied. A SAS data set can be opened concurrently by more than one SAS session or by more than one statement, window, or procedure within a single session. By default, SAS procedures allow the greatest degree of concurrent access possible, and they guarantee the integrity of the data and the data analysis. Therefore, you typically use the CNTLLEV= data set option in these situations:

- when your application controls the access to the data, such as in SAS Component Language (SCL), SAS/IML software, or DATA step programming
- when you access data through an interface engine that does not provide member-level control of the data.

If you use CNTLLEV=REC and the SAS procedure needs member-level control for integrity of the data analysis, SAS prints a warning to the SAS log. The warning states that inaccurate or unpredictable results can occur if the data is updated by another process during the analysis.

Example: Changing the Shared Access Level

In this example, the first SET statement includes the CNTLLEV= data set option in order to override the default level of shared access from member-level control to record-level control. The second SET statement opens the SAS data set with the default member-level control.

```

libname mylib 'c:\mydata';

data mylib.new;
set mylib.fuel(cntllev=rec) point=obsnum;
/* ... more SAS code ... */
set mylib.fuel;
by area;
run;

```

COMPRESS= Data Set Option

Specifies how observations are compressed in a new output SAS data set.

Valid in:	DATA step and PROC steps
Categories:	Data Set Control CAS
Restrictions:	When using this option with the CAS server, use YES and NO values only. Use with output data sets only.

Syntax

COMPRESS=[NO](#) | [YES](#) | [CHAR](#) | [BINARY](#)

Syntax Description

NO

specifies that the observations in a newly created SAS data set are uncompressed (maintaining fixed-length records).

YES | CHAR

specifies that the observations in a newly created SAS data set are compressed (producing variable-length records) by using RLE (Run Length Encoding). RLE compresses observations by reducing repeated runs of the same character (including blanks) to 2-byte or 3-byte representations.

Alias ON

BINARY

specifies that the observations in a newly created SAS data set are compressed (producing variable-length records) by using RDC (Ross Data Compression). RDC combines run-length encoding and sliding-window compression to compress the file by representing repeated byte patterns more efficiently.

Note: This method is highly effective for compressing medium to large (several hundred bytes or larger) blocks of binary data (character and numeric variables). Because the compression function operates on a single record at a time, the

record length needs to be several hundred bytes or larger for effective compression.

.....

Details

When copying compressed between V9 engine libraries and CAS engines libraries that use the COPY procedure or PROC DATASETS COPY statement, the following actions occur:

- If a SAS data set is compressed, it retains the COMPRESS=YES value on a CAS table.
- If a CAS table is compressed, it converts to a SAS data set with the COMPRESS=CHAR value.

Compressing a file reduces the number of bytes that are required to represent each observation. Advantages of compressing a file include reduced storage requirements for the file and fewer I/O operations to read or write to the data during processing. However, more CPU resources are required to read a compressed file (because of the overhead of uncompressing each observation). There are situations where the resulting file size might increase rather than decrease.

Use the COMPRESS= data set option to compress an individual file. Specify the option for output data sets only. That is, specify data sets named in the DATA statement of a DATA step or in the OUT= option of a SAS procedure. Use the COMPRESS= data set option only when you are creating a SAS data file (member type DATA). You cannot compress SAS views, because they contain no data. The COPY procedure does not support data set options. Therefore, you cannot use the COMPRESS= data set option in PROC COPY or a COPY statement from PROC DATASETS.

TIP When you use the V9 engine to compress an OUTPUT data set that is generated by PROC COPY, you can use the COMPRESS=YES system option before the PROC COPY statement with the NOCLONE option.

```
options compress=yes;
proc copy in=work out=new noclone;
select x;
run;
```

The CAS engine does not support the COMPRESS= system option. However, you can use the COMPRESS LIBNAME option on the output library, and then use the PROC COPY statement with the NOCLONE option.

After a file is compressed, the setting is a permanent attribute of the file. To change the setting, you must re-create the file. That is, to uncompress a file, specify COMPRESS=NO for a DATA step that copies the compressed file.

In general, COMPRESS=CHAR provides good compression when single bytes repeat; COMPRESS=BINARY provides good compression when strings of bytes

repeat. It is more costly to look for strings of bytes that repeat, than to look for single bytes that repeat. For examples, see [“Example 1: COMPRESS=CHAR” on page 18](#) and [“Example 2: COMPRESS=BINARY” on page 19](#).

Comparisons

- The COMPRESS= data set option overrides the COMPRESS= option in the LIBNAME statement and the COMPRESS= system option.
- If the COMPRESS= data set option or LIBNAME option is not set, then the value of the COMPRESS= system option is used. The COMPRESS= system option default value is NO.
- The data set option POINTOBS=YES, which is the default, determines that a compressed data set can be processed with random access (by observation number) rather than sequential access. With random access, you can specify an observation number in the FSEDIT procedure and the POINT= option in the SET and MODIFY statements.
- When you create a compressed file, you can also specify REUSE=YES (as a data set option or system option) to track and reuse space. With REUSE=YES, new observations are inserted in available space when other observations are updated or deleted. When the default REUSE=NO is in effect, new observations are appended to the existing file.
- POINTOBS=YES and REUSE=YES are mutually exclusive. That is, they cannot be used together. REUSE=YES takes precedence over POINTOBS=YES. If you set REUSE=YES, SAS automatically sets POINTOBS=NO.
- The TAPE engine supports the COMPRESS= data set option, but the engine does not support the COMPRESS= system option.
- The XPORT engine does not support compression.

Examples

Example 1: COMPRESS=CHAR

```
libname mylib 'c:\mydata';

data mylib.CharRepeats (compress=char);
  length ca $ 200;
  do i=1 to 100000;
    ca='aaaaaaaaaaaaaaaaaaaaa';
    cb='bbbbbbbbbbbbbbbbbbbbbb';
    cc='cccccccccccccccccccc';
    output;
  end;
run;
```

This message is written to the SAS log:

NOTE: Compressing data set MYLIB.CHARREPEATS decreased size by 88.55 percent.

Compressed is 45 pages; uncompressed would require 393 pages.

Example 2: COMPRESS=BINARY

```
libname mylib 'c:\mydata';

data mylib.StringRepeats(compress=binary);
  length cabcd $ 200;
  do i=1 to 1000000;
    cabcd='abcdabcdabcdabcdabcdabcdabcdabcd';
    cefgh='efghefghefghefghefghefghefghefghefgh';
    cijkl='ijklijklijklijklijklijklijklijkl';
    output;
  end;
run;
```

This message is written to the SAS log:

NOTE: Compressing data set MYLIB.STRINGREPEATS decreased size by 70.27 percent.
Compressed is 1239 pages; uncompressed would require 4167 pages.

See Also

- [“Compression in SAS” in SAS Programmer’s Guide: Essentials](#)

Data Set Options:

- [“POINTOBS= Data Set Option” on page 64](#)
- [“REUSE= Data Set Option” on page 77](#)

Statements:

- [“LIBNAME Statement” in SAS Global Statements: Reference](#)

System Options:

- [“COMPRESS= System Option” in SAS System Options: Reference](#)
- [“REUSE= System Option” in SAS System Options: Reference](#)

DLDMGACTION= Data Set Option

Specifies the action to take when a SAS data set in a SAS library is detected as damaged.

Valid in: DATA step and PROC steps

Category: Data Set Control

Defaults: For Windows and UNIX, the shipped default is REPAIR for interactive mode and FAIL for batch mode.

For z/OS, the shipped default is PROMPT for interactive mode and REPAIR for batch mode.

Restriction: Use in the Compute Server, but not in CAS.

Syntax

DLDMGACTION=[FAIL](#) | [ABORT](#) | [REPAIR](#) | [NOINDEX](#) | [PROMPT](#)

Syntax Description

FAIL

stops the step and issues an error message to the log immediately. This is the default for batch mode.

ABORT

stops the step, issues an error message to the log, and terminates the SAS session.

REPAIR

attempts to automatically repair a damaged data set on the next attempt to open the damaged file. The data set might truncate at the point of damage. The REPAIR option re-creates the index or indexes. If the damage is too severe, the auto repair attempt might not be successful.

NOINDEX

automatically repairs the data file without the indexes and integrity constraints. The repair also deletes the index file and updates the data file to reflect the disabled indexes and integrity constraints. The repair limits the data file to be opened only in INPUT mode. A warning is written to the SAS log instructing you to execute the PROC DATASETS REBUILD statement to correct or delete the disabled indexes and integrity constraints.

See [“DLDMGACTION= System Option” in SAS System Options: Reference](#)

PROMPT

displays a dialog box that asks you to select the FAIL, ABORT, REPAIR, or NOINDEX action.

DROP= Data Set Option

For an input data set, excludes the specified variables from processing; for an output data set, excludes the specified variables from being written to the data set.

Valid in: DATA step and PROC steps

Categories: Variable Control
CAS

Syntax

DROP=*variable(s)*

Syntax Description

variable(s)

lists one or more variable names. You can list the variables in any form that SAS allows.

Details

If the option is associated with an input data set, the variables are not available for processing. If the DROP= data set option is associated with an output data set, SAS does not write the variables to the output data set, but they are available for processing.

Comparisons

The DROP= data set option differs from the DROP statement in these ways:

- In DATA steps, the DROP= data set option can apply to input and output data sets. The DROP statement applies only to output data sets.
- In DATA steps, when you create multiple output data sets, use the DROP= data set option to write different variables to different data sets. The DROP statement applies to all output data sets.
- In PROC steps, you can use only the DROP= data set option, not the DROP statement.

Examples

Example 1: Excluding Variables from Input

In this example, the variables SALARY and GENDER are not included in processing, and they are not written to either output data set:

```
libname mylib 'c:\mydata';

data plan1 plan2;
  set payroll(drop=salary gender);
  if hired<'01jan98'd then output plan1;
  else output plan2;
run;
```

You cannot use SALARY or GENDER in any logic in the DATA step because DROP= prevents the SET statement from reading SALARY and GENDER from PAYROLL.

Example 2: Processing Variables without Writing Them to a Data Set

In this example, SALARY and GENDER are not written to PLAN2, but they are written to PLAN1:

```
libname mylib 'c:\mydata';

data plan1 plan2(drop=salary gender);
  set payroll;
  if hired<'01jan98'd then output plan1;
  else output plan2;
run;
```

See Also

Data Set Options:

- [“KEEP= Data Set Option” on page 45](#)

Statements:

- [“DROP Statement” in SAS DATA Step Statements: Reference](#)

ENCRYPT= Data Set Option

Specifies whether to encrypt an output SAS data set.

Valid in:	DATA step and PROC steps
Category:	Data Set Control
Default:	ENCRYPT=NO
Restrictions:	Use in the Compute Server, but not in CAS. Use with output data sets only. If you copy an AES2 data set to an SPD Engine data set, the encryption is changed to AES. SAS 9.4M5 or later software is required to use AES2 encryption.

Syntax

ENCRYPT=[AES](#) | [AES2](#) | [NO](#) | [YES](#)

Syntax Description

AES or AES2

encrypts the file by using the AES (Advanced Encryption Standard) algorithm. AES provides enhanced encryption by using SAS/SECURE software, which is included with Base SAS software. AES2 specifies to use a stronger key generation algorithm. You must specify the ENCRYPTKEY= data set option when you are using ENCRYPT=AES or ENCRYPT=AES2, or you must have a recorded key in the metadata-bound library where the data set is created. For more information, see [“ENCRYPTKEY= Data Set Option” on page 27](#).

Restrictions The tape engine does not support ENCRYPT=AES or ENCRYPT=AES2. Use ENCRYPT=YES for tape engine encryption.

The SPD Engine does not support ENCRYPT=AES2. If a default Base SAS data set with AES2 encryption is copied to create a new SPD Engine data set, the encryption converts to AES. A warning is written to the log.

CAUTION **Record all ENCRYPTKEY= values when you are using ENCRYPT=AES or ENCRYPT=AES2.** If you forget to record the ENCRYPTKEY= value, you lose your data. SAS cannot assist you in recovering the ENCRYPTKEY= value. The following note is written to the log:

Note: If you lose or forget the ENCRYPTKEY= value, there will be no way to open the file or recover the data.

NO

does not encrypt the file.

YES

encrypts the file by using the SAS Proprietary algorithm. This encryption uses passwords that are stored in the data set. At a minimum, you must specify the READ= data set option or the PW= data set option at the same time that you specify ENCRYPT=YES. Because the encryption method uses passwords, you cannot change any password on an encrypted data set without re-creating the data set.

CAUTION

Record all passwords when you are using ENCRYPT=YES. If you forget the passwords, you cannot reset them without assistance from SAS. This is a time-consuming and resource-intensive process.

Details

To use ENCRYPT=YES data files, you must have SAS 6.11 or later. When you use ENCRYPT=YES, these rules apply:

- To copy an encrypted data file, the output engine must support the encryption.
- If the data file is encrypted, all associated indexes are also encrypted.

- You cannot use PROC CPORT on SASProprietary encrypted data files.

Note: SAS views do not contain data. Therefore, encryption is not necessary.

To use encrypted AES or AES2 data files, you must use SAS 9.4 or later and SAS/SECURE software. In addition, when you use ENCRYPT=AES or ENCRYPT=AES2, these rules apply:

- You must use the ENCRYPTKEY= data set option when creating a data set with AES or AES2 encryption, or you must have a recorded encrypt key value in the metadata-bound library where the data set was created.
- To copy an encrypted AES data file, the output engine must support AES encryption.
- In Base SAS, data files with referential integrity constraints can use AES encryption. All primary key and foreign key data files must use the same encryption key that opens all referencing foreign key and primary key data files.

You cannot change the ENCRYPTKEY= value on an AES-encrypted data file without re-creating the data file or using the AUTHLIB procedure to change the recorded key of a metadata-bound library. For more information, see [“AUTHLIB Procedure” in Base SAS Procedures Guide](#).

Note: Encryption requires approximately the same amount of CPU resources as compression.

Examples

Example 1: Using the ENCRYPT=YES Option

The following example encrypts the data set by using the SAS Proprietary algorithm:

```
libname mylib 'c:\mydata';

data mylib.salary(encrypt=yes read=green);
  input name $ yrsal bonuspct;
  datalines;
Muriel      34567  3.2
Bjorn       74644  2.5
Freda       38755  4.1
Benny       29855  3.5
Agnetha     70998  4.1
;
```

To use this data set, specify the READ= password:

```
proc contents data=mylib.salary(read=green);
run;
```

Output 2.1 ENCRYPT=YES

The SAS System			
The CONTENTS Procedure			
Data Set Name	MYLIB.SALARY	Observations	5
Member Type	DATA	Variables	3
Engine	V9	Indexes	0
Created	04/30/2014 13:37:05	Observation Length	24
Last Modified	04/30/2014 13:37:05	Deleted Observations	0
Protection	READ	Compressed	NO
Data Set Type		Sorted	NO
Encrypted	YES		
Label			
Data Representation	WINDOWS_64		
Encoding	wlatin1 Western (Windows)		

Engine/Host Dependent Information	
Data Set Page Size	65536
Number of Data Set Pages	1
First Data Page	1
Max Obs per Page	2715
Obs in First Data Page	5
Number of Data Set Repairs	0
ExtendObsCounter	YES
Filename	c:\mylib\salary.sas7bdat
Release Created	9.0401MD
Host Created	X64_7PRO

Example 2: Using the ENCRYPT=AES Option

The following example encrypts the data set by using the AES algorithm:

```
libname mylib 'c:\mydata';

data mylib.salary(encrypt=aes encryptkey=green);
  input name $ yrsal bonuspct;
  datalines;
Muriel      34567   3.2
Bjorn       74644   2.5
```

```

Freda      38755  4.1
Benny      29855  3.5
Agnetha    70998  4.1
;

```

To use this data set, specify the ENCRYPTKEY= key value:

```

proc contents data=mylib.salary(encryptkey=green);
run;

```

Output 2.2 ENCRYPT=AES

The SAS System

The CONTENTS Procedure

Data Set Name	MYLIB.SALARY	Observations	5
Member Type	DATA	Variables	3
Engine	V9	Indexes	0
Created	04/30/2014 13:42:52	Observation Length	24
Last Modified	04/30/2014 13:42:52	Deleted Observations	0
Protection		Compressed	NO
Data Set Type		Sorted	NO
Encrypted	AES		
Label			
Data Representation	WINDOWS_64		
Encoding	wlatin1 Western (Windows)		

Engine/Host Dependent Information

Data Set Page Size	65536
Number of Data Set Pages	1
First Data Page	1
Max Obs per Page	2715
Obs in First Data Page	5
Number of Data Set Repairs	0
ExtendObsCounter	YES
Filename	c:\mylib\salary.sas7bdat
Release Created	9.0401M0
Host Created	X64_7PRO

See Also

- [“SAS Data File Encryption” in SAS Programmer’s Guide: Essentials](#)

Data Set Options:

- [“ALTER= Data Set Option” on page 9](#)
- [“PW= Data Set Option” on page 66](#)
- [“READ= Data Set Option” on page 68](#)
- [“WRITE= Data Set Option” on page 98](#)

ENCRYPTKEY= Data Set Option

Specifies a key value for AES (Advanced Encryption Standard) encryption.

Valid in:	DATA step and PROC steps
Category:	Data Set Control
Range:	1 to 64 bytes
Restrictions:	Use in the Compute Server, but not in CAS. Use only with SAS 9.4 or later. Use only with AES-encrypted data files.
Note:	Check your log after this operation to ensure encryption key security. For more information, see “Blotting Passwords and Encryption Key Values” in SAS Programmer’s Guide: Essentials .

Syntax

ENCRYPTKEY=*key-value*

Syntax Description

key-value

assigns an encrypt key value. You must specify the ENCRYPTKEY= data set option when you are using ENCRYPT=AES or ENCRYPT=AES2 unless you are creating a data set in a metadata-bound library with a recorded key. The key value can be up to 64 bytes long. To create an ENCRYPTKEY= key value with or without quotation marks, follow these rules:

No quotation marks:

- use alphanumeric characters and underscores only
- can be up to 64 bytes long
- use uppercase and lowercase letters

- must start with a letter
- cannot include blank spaces
- is not case sensitive

```
%let mykey=abcdefghi12;
encryptkey=&mykey
encryptkey=key_value
encryptkey=key_value1
```

Single quotation marks:

- use alphanumeric, special, and DBCS characters
- can be up to 64 bytes long
- use uppercase and lowercase letters
- can include blank spaces, but cannot contain all blanks
- is case sensitive

```
encryptkey='key_value'
encryptkey='1234*##mykey'
```

Double quotation marks:

- use alphanumeric, special, and DBCS characters
- can be up to 64 bytes long
- use uppercase and lowercase letters
- can include blank spaces, but cannot contain all blanks
- is case sensitive

```
encryptkey="key_value"
encryptkey="1234*##mykey"
%let mykey=Abcdefghi12;
encryptkey="&mykey"
```

When the ENCRYPTKEY= key value uses DBCS characters, the 64-byte limit applies to the character string after it has been transcoded to UTF-8 encoding. You can use the following DATA step to calculate the length in bytes of a key value in DBCS:

```
data _null_;
  key=length(unicode('key-value','UTF8'));
  put 'key length=' key;
run;
```

Interaction You cannot change the key value on an AES-encrypted data set without re-creating the data set or using the AUTHLIB procedure to change the recorded key in a metadata-bound library. For more information, see [“AUTHLIB Procedure” in Base SAS Procedures Guide](#).

Details

CAUTION

Record all ENCRYPTKEY= values when you are using ENCRYPT=AES or ENCRYPT=AES2. If you forget to record the ENCRYPTKEY= value, you lose your data. SAS cannot assist you in recovering the ENCRYPTKEY= value.

You must use the ENCRYPTKEY= option when you are creating or accessing a SAS data set with AES encryption.

The ENCRYPTKEY= data set option does not protect the file from deletion or replacement. Encrypted data sets can be deleted using any of the following scenarios without specifying an ENCRYPTKEY= key value:

- the KILL option in PROC DATASETS
- the DROP statement in PROC SQL
- the DELETE statement in PROC DATASETS or the DELETE procedure

The ENCRYPTKEY= option prevents access to the contents of the file only. To protect the file from deletion or replacement, the file must also contain an ALTER= or PW= password.

You must specify the ENCRYPTKEY= key value when you copy AES-encrypted data files. The value follows the data set name in the SELECT statement. The following example uses the SELECT statement:

```
copy in=OldLib out=NewLib;
    select salary(encryptkey=key-value);
run;
```

When you are working with data files that are protected with AES encryption that do not have a recorded encryption key in a metadata-bound library where the data set resides, you can specify the value in the AGE, APPEND, AUDIT, CONTENTS, MODIFY, REBUILD, and REPAIR statements in the DATASETS procedure. You must also specify the value when the CHANGE statement refers to a specific generation data set by using a relative reference to the value:

```
change OldName(gennum=-1 encryptkey=key-value)=NewName;
run;
```

The option can be specified either in parentheses after the name of the SAS data file or after a forward slash.

CAUTION

When you are using referential integrity constraints, all primary key and foreign key data files that reference each other must use the same encryption key. For more information, see [“AES Encryption and Referential Integrity Constraints” in SAS Programmer’s Guide: Essentials](#).

You can use a macro variable as the ENCRYPTKEY= key value. The following code defines a macro variable:

```
%let secret=myvalue;
```

The following code uses the macro variable as the ENCRYPTKEY= key value:

```
data my.dsname(encrypt=aes encryptkey="&secret");
```

When you specify a macro variable for the ENCRYPTKEY= key value, you must enclose the macro variable in double quotation marks. If you do not use double quotation marks, unpredictable results can occur.

Example: Using the ENCRYPTKEY= Option

This example uses the ENCRYPT=AES option:

```
libname mylib 'c:\mydata';

data salary(encrypt=aes encryptkey=green);
  input name $ yrsal bonuspct;
  datalines;
Muriel      34567  3.2
Bjorn       74644  2.5
Freda       38755  4.1
Benny       29855  3.5
Agnetha     70998  4.1
;
```

To use this data set, specify the ENCRYPTKEY= key value:

```
proc contents data=salary(encryptkey=green);
run;
```

See Also

[“SAS Data File Encryption” in SAS Programmer’s Guide: Essentials](#)

EXTENDOBSCOUNTER= Data Set Option

Specifies whether to extend the maximum observation count in a new output SAS data file.

Valid in:	DATA step and PROC steps
Category:	Data Set Control
Alias:	EOC=
Defaults:	For Windows and z/OS, the default is YES. For UNIX, the default is not set.
Restrictions:	Use in the Compute Server, but not in CAS. Use with output data files only. Use with the V9 engine only. If you copy a file, the extended observation count attribute is not inherited.

Syntax

EXTENDOBSCOUNTER=YES | NO

Syntax Description

YES

specifies to use the default maximum observation count of $2^{63}-1$ or approximately 9.2 quintillion observations in a newly created SAS data file.

Restriction A SAS data file that has the extended observation count attribute cannot be used by releases prior to SAS 9.3. To remove the extended observation count attribute, the file must be re-created. If you plan to use the file in SAS 9.2 or earlier releases, then set EXTENDOBSCOUNTER=NO when you create the file.

Interaction EXTENDOBSCOUNTER=YES is ignored if the data representation is a 64-bit UNIX operating environment.

NO

is a compatibility setting for a newly created SAS data file, to enable its use in releases prior to SAS 9.3.

Details

The behavior depends on your operating environment:

- Under UNIX environments, by default the EXTENDOBSCOUNTER= option is not set. The extended observation count feature is not necessary under 64-bit UNIX. However, if you specify the OUTREP= option, and the data representation is not a 64-bit UNIX operating environment, then SAS automatically sets EXTENDOBSCOUNTER=YES. SAS adds the extended observation count feature for compatibility with environments other than UNIX where it might be necessary. (Many customers specify OUTREP= when they create a table for use in a different environment. This practice can help you avoid the limitations of CEDA processing.)
- Under Windows and z/OS, by default EXTENDOBSCOUNTER=YES. Files are created with the enhanced file format and the extended observation count attribute.

Note that backward compatibility is an issue only if the file is used in SAS 9.2 or earlier releases:

- Under UNIX, if you specify OUTREP= and plan to use the file in SAS 9.2 or earlier releases, specify EXTENDOBSCOUNTER=NO. If you do not specify OUTREP=, then you do not need to specify EXTENDOBSCOUNTER=NO.
- Under Windows or z/OS, if you plan to use the file in SAS 9.2 or earlier releases, specify EXTENDOBSCOUNTER=NO.

See Also

Statements:

- [“EXTENDOBSCOUNTER=YES | NO” in SAS Global Statements: Reference](#)

System Options:

- [“EXTENDOBSCOUNTER= System Option” in SAS System Options: Reference](#)

FILECLOSE= Data Set Option

Specifies how a tape is positioned when a SAS data set is closed.

Valid in: DATA step and PROC steps

Category: Miscellaneous

Restriction: Use in the Compute Server, but not in CAS.

See: FILECLOSE= Data Set Option in the documentation for the UNIX operating environment.

CAUTION: **The option values are not recognized by all operating environments.** Additional values are available on some operating environments. For more information about using SAS libraries that are stored on tape, see the SAS documentation for your operating environment.

Syntax

FILECLOSE=[DISP](#) | [LEAVE](#) | [REREAD](#) | [REWIND](#)

Syntax Description

DISP

positions the tape volume according to the disposition that is specified in the operating environment's control language.

LEAVE

positions the tape at the end of the file that was recently processed. Use FILECLOSE=LEAVE if you are not repeatedly accessing the same files in a SAS program, but you are accessing one or more subsequent SAS files on the same tape.

REREAD

positions the tape volume at the beginning of the file that was recently processed. Use FILECLOSE=REREAD if you are accessing the same SAS data set on tape several times in a SAS program.

REWIND

rewinds the tape volume to the beginning. Use FILECLOSE=REWIND if you are accessing one or more previous SAS files on the same tape, but you are not repeatedly accessing the same files in a SAS program.

FIRSTOBS= Data Set Option

Specifies the first observation that SAS processes in a SAS data set.

Valid in:	DATA step and PROC steps
Category:	Observation Control
Restrictions:	Use in the Compute Server, but not in CAS. Valid for input (read) processing only. Cannot use with PROC SQL views.

Syntax

FIRSTOBS=*n* | *nK* | *nM* | *nG* | *hexX* | MIN | MAX

Syntax Description

n* | *nK* | *nM* | *nG

specifies the number of the first observation to process in multiples of 1 (bytes); 1,024 (kilobytes); 1,048,576 (megabytes); or 1,073,741,824 (gigabytes). For example, a value of 8 specifies the eighth observation, and a value of 3k specifies 3,072.

Note: You can also use the KB, MB, and GB syntax notations.

hexX

specifies the number of the first observation to process as a hexadecimal value. You must specify the value beginning with a number (0–9), followed by an X. For example, the value 2dx sets the 45th observation as the first observation to process.

MIN

sets the number of the first observation to process to 1. This is the default.

MAX

sets the number of the first observation to process to the maximum number of observations in the data set. This number can be up to the largest eight-byte, signed integer, which is $2^{63}-1$, or approximately 9.2 quintillion observations.

Details

The FIRSTOBS= data set option affects a single, existing SAS data set. Use the FIRSTOBS= system option to affect all steps for the duration of your current SAS session.

FIRSTOBS= is valid for input (read) processing only. Specifying FIRSTOBS= is not valid for output or update processing.

You can apply FIRSTOBS= processing to WHERE processing.

If you delete data set observations by using the VIEWTABLE window, the FIRSTOBS= option gives incorrect results if the value assigned to the FIRSTOBS= option is greater than the number of the observation that were deleted. This behavior occurs because deleting data set observations by using the VIEWTABLE window only flags the observation for deletion. The observation is not physically removed from the data set; it makes the observation unusable. For more information, see [“Working with VIEWTABLE” in SAS Programmer’s Guide: Essentials](#).

Comparisons

- The FIRSTOBS= data set option overrides the FIRSTOBS= system option for the individual data set.
- When the FIRSTOBS= data set option specifies a starting point for processing, the OBS= data set option specifies an ending point. The two options are often used together to define a range of observations to be processed.
- When external files are read, the FIRSTOBS= option in the INFILE statement specifies which record to read first.

Example: Using the FIRSTOBS= Data Set Option

This PROC step prints the data set STUDY, beginning with observation 20:

```
proc print data=study(firstobs=20);  
run;
```

This SET statement uses FIRSTOBS= and OBS= to read only observations 5 through 10 from the data set STUDY. The data set NEW contains six observations.

```
data new;  
  set study(firstobs=5 obs=10);  
run;  
  
proc print data=new;  
run;
```

See Also

Data Set Options:

- [“OBS= Data Set Option” on page 52](#)

Statements:

- [“INFILE Statement” in SAS DATA Step Statements: Reference](#)
- [“WHERE Statement” in SAS DATA Step Statements: Reference](#)

System Options:

- [“FIRSTOBS= System Option” in SAS System Options: Reference](#)

GENMAX= Data Set Option

Requests generations for a new data set, modifies the number of generations for an existing data set, and specifies the maximum number of versions.

Valid in:	DATA step and PROC steps
Category:	Data Set Control
Default:	0
Restrictions:	Use in the Compute Server, but not in CAS. Use with output data sets only.

Syntax

GENMAX=*number-of-generations*

Syntax Description

number-of-generations

requests generations for a data set and specifies the maximum number of versions to maintain. The value can be from 0 to 1,000. The default is GENMAX=0, which means that no generation data sets are requested.

Details

You use GENMAX= to request generations for a new data set and to modify the number of generations for an existing data set. The first time the data set is replaced, SAS keeps the replaced version and appends a 4-character version

number to its member name. The member name includes # and a three-digit number. For example, for a data set named A, a historical version would be A#001.

After generations of a data set are requested, the member name is limited to 28 characters (rather than 32). The last four characters are reserved for the appended version number. When the GENMAX= data set option is set to 0, the member name can be up to 32 characters.

If you reduce the number of generations for an existing data set, SAS deletes the oldest versions above the new limit.

Examples

Example 1: Requesting Generations When You Create a Data Set

The DATA step creates a data set named Work.A that can have as many as 10 generations (one current version and nine historical versions):

```
libname mylib 'c:\mydata';

data a(genmax=10);
    x=1;
    output;
run;
```

Example 2: Modifying the Number of Generations on an Existing Data Set

The number of generations on the data set MYLIB.A is changed to 4:

```
proc datasets lib=mylib;
    modify a(genmax=4);
run;
quit;
```

See Also

[“GENNUM= Data Set Option” on page 36](#)

GENNUM= Data Set Option

Specifies a particular generation of a SAS data set.

Valid in: DATA step and PROC steps

Category: Data Set Control

Restrictions: Use in the Compute Server, but not in CAS.

Use with input data sets only.

Syntax

GENNUM=*integer*

Syntax Description

integer

is a number that references a specific version from a generation group. Specifying a positive number is an absolute reference to a specific generation number that is appended to the name of a data set. Specifying a negative number is a relative reference to a historical version in relation to the base (current) version, from the youngest to the oldest. Typically, a value of 0 refers to the base version.

The DATASETS procedure provides a variety of statements for which specifying GENNUM= has additional functionality:

- For the DATASETS and DELETE statements, GENNUM= supports the additional values ALL, HIST, and REVERT.
- For the DELETE procedure, GENNUM= supports the additional values ALL, HIST, and REVERT.
- For the CHANGE statement, GENNUM= supports the additional value ALL.
- For the CHANGE statement, specifying GENNUM=0 refers to all versions rather than the base version only.

Details

Note: Do not confuse the GENNUM variable value in the CONTENTSOUT= data set with the GEN variable value from DICTIONARY tables.

GENNUM from a CONTENTS procedure or statement refers to a specific generation of a data set. GEN from DICTIONARY tables refers to the total number of generations for a data set.

Each SAS procedure is designed and configured to deliver specific information. The output from PROC CONTENTS and DICTIONARY.TABLES is not interchangeable.

After generations for a data set have been requested using the GENMAX= data set option, use GENNUM= to request a specific version. For example, specifying GENNUM=3 refers to the historical version #003; specifying GENNUM= -1 refers to the youngest historical version.

After 999 replacements, the youngest version would be #999. After 1,000 replacements, SAS rolls over the youngest version number to #000. Therefore, if you want the historical version #000, specify GENNUM=1000.

Both an absolute reference and a relative reference refer to a specific version. A relative reference does not skip deleted versions. Therefore, when working with a generation group that includes one or more deleted versions, using a relative reference results in an error if the version that is being referenced has been deleted. For example, if you have the base version AIR and three historical versions (AIR#001, AIR#002, and AIR#003), delete AIR#002. The following statements return an error because AIR#002 does not exist.

```
proc print data=air (gennum= -2);
run;
```

Examples

Example 1: Requesting a Version Using an Absolute Reference

This example prints the historical version #003 for data set A, using an absolute reference:

```
proc print data=a(gennum=3);
run;
```

Example 2: Requesting a Version Using a Relative Reference

The following PRINT procedure prints the data set three versions back from the base version:

```
proc print data=a(gennum=-3);
run;
```

See Also

- [“DATASETS Procedure” in Base SAS Procedures Guide](#)
- [“GENMAX= Data Set Option” on page 35](#)

IDXNAME= Data Set Option

Directs SAS to use a specific index to match the conditions of a WHERE expression.

Valid in:	DATA step and PROC steps
Category:	User Control of SAS Index Usage
Restrictions:	Use in the Compute Server, but not in CAS. Use with input data sets only. Mutually exclusive with IDXWHERE= data set option

Syntax

IDXNAME=*index-name*

Syntax Description

index-name

specifies the name (up to 32 characters) of a simple or composite index for the SAS data set. SAS does not attempt to determine whether the specified index is the best one or whether a sequential search might be more resource efficient.

Interaction The specification is not a permanent attribute of the data set and is valid only for the current use of the data set.

Tip To request that IDXNAME= usage be noted in the SAS log, specify the system option MSGLEVEL=I.

Details

To satisfy the conditions of a WHERE expression for an indexed SAS data set, SAS identifies zero or more candidate indexes that could be used to optimize the WHERE expression. From the list of candidate indexes, SAS determines the following:

- the candidate index that provides the best performance
- the rejection of all the indexes if a sequential pass of the data is more efficient

Because the index SAS selected cannot always provide the best optimization, you can direct SAS to use one of the candidate indexes by specifying the IDXNAME= data set option. If you specify an index that SAS does not identify as a candidate index, then IDXNAME= does not process the request. That is, IDXNAME= does not enable you to specify an index that would produce incorrect results.

Comparisons

- **IDXWHERE=** enables you to override the SAS decision about whether to use an index, whereas **INDEXNAME=** enables you to direct SAS to use a specific index.

Example: Specifying an Index

This example uses the IDXNAME= data set option in order to direct SAS to use a specific index to optimize the WHERE expression. SAS then disregards the possibility that a sequential search of the data set might be more resource efficient. SAS does not attempt to determine whether the specified index is the best one. (The EMPNUM index was not created with the NOMISS option.)

```
libname mylib 'c:\mydata';

data mylib.empnew;
  set mylib.employee (idxname=empnum);
  where empnum < 2000;
run;
```

See Also

- [“Using Indexes” in SAS Programmer’s Guide: Essentials](#)

Data Set Option:

- [“IDXWHERE= Data Set Option” on page 40](#)

IDXWHERE= Data Set Option

Specifies whether SAS uses an index search or a sequential search to match the conditions of a WHERE expression.

Valid in:	DATA step and PROC steps
Category:	User Control of SAS Index Usage
Restrictions:	Use in the Compute Server, but not in CAS. Use with input data sets only. Mutually exclusive with IDXNAME= data set option

Syntax

IDXWHERE=[YES](#) | [NO](#)

Syntax Description

YES

tells SAS to choose the best index to optimize a WHERE expression, and to disregard the possibility that a sequential search of the data set might be more resource efficient.

NO

tells SAS to ignore all indexes and satisfy the conditions of a WHERE expression with a sequential search of the data set.

Note: You cannot use IDXWHERE= to override the use of an index to process a BY statement.

Details

By default, to satisfy the conditions of a WHERE expression for an indexed SAS data set, SAS decides whether to use an index or to read the data set sequentially. The software estimates the relative efficiency and chooses the method that is more efficient.

You might need to override the software's decision by specifying the IDXWHERE= data set option. The decision is based on general rules that might not always produce the best results. By specifying the IDXWHERE= data set option, you can determine the processing method.

Note: The specification is not a permanent attribute of the data set and is valid only for the current use of the data set.

Note: If you issue the system option MSGLEVEL=I, you can request that IDXWHERE= usage be noted in the SAS log if the setting affects index processing.

Comparisons

- IDXNAME= enables you to direct SAS to use a specific index, whereas INDEXWHERE= enables you to override the SAS decision about whether to use an index.

Examples

Example 1: Specifying Index Usage

This example uses the IDXWHERE= data set option to tell SAS to decide which index is the best to optimize the WHERE expression. SAS then disregards the possibility that a sequential search of the data set might be more resource efficient:

```
libname mylib 'c:\mydata';

data mylib.empnew;
    set mylib.employee (idxwhere=yes);
    where empnum < 2000;
run;
```

Example 2: Specifying No Index Usage

This example uses the IDXWHERE= data set option to perform the following tasks:

- tell SAS to ignore any index
- satisfy the conditions of the WHERE expression with a sequential search of the data set

```
libname mylib 'c:\mydata';

data mylib.empnew;
    set mylib.employee (idxwhere=no);
    where empnum < 2000;
run;
```

See Also

- [“Using Indexes” in SAS Programmer’s Guide: Essentials](#)
- [WHERE-expression Processing](#)

Data Set Option:

- [“IDXNAME= Data Set Option” on page 38](#)

INDEX= Data Set Option

Defines an index for a new output SAS data set.

Valid in:	DATA step and PROC steps
Category:	Data Set Control
Restrictions:	Use in the Compute Server, but not in CAS. Use with output data sets only.

Syntax

INDEX=(*index-specification* <*index-specification-2* ...>)

Syntax Description

index-specification

names and describes a simple or a composite index to be built. *Index-specification* has this form:

index= (*variable(s)* </UNIQUE> </NOMISS>)

index

is the name of a variable that forms the index or the name that you choose for a composite index.

variable(s)

is a list of variables to use in building a composite index.

UNIQUE

specifies that the values of the key variables must be unique. If you specify UNIQUE for a new data set and multiple observations have the same values for the index variables, the index is not built. A slash (/) must precede the UNIQUE option.

NOMISS

excludes all observations with missing values from the index. Observations with missing values are still read from the data set, but not through the index. A slash (/) must precede the NOMISS option.

Examples

Example 1: Defining a Simple Index

This INDEX= data set option defines a simple index for the SSN variable:

```
data new(index=(ssn));
/* ... more SAS code ... */
run;
```

Example 2: Defining a Composite Index

This INDEX= data set option defines a composite index named CITYST that uses the CITY and STATE variables:

```
data new(index=(cityst=(city state)));
/* ... more SAS code ... */
run;
```

Example 3: Defining a Simple and a Composite Index

This INDEX= data set option defines a simple index for SSN and a composite index for CITY and STATE:

```
data new(index=(ssn cityst=(city state)));
/* ... more SAS code ... */
run;
```

Example 4: Defining a Simple Index with the UNIQUE Option

This INDEX= data set option defines a simple index for the SSN variable with unique values:

```
data new(index=(ssn /unique));
/* ... more SAS code ... */
run;
```

Example 5: Defining a Simple Index with the NOMISS Option

This INDEX= data set option defines a simple index for the SSN variable, excluding all observations with missing values from the index:

```
data new(index=(ssn /nomiss));
```

```
/* ... more SAS code ... */
run;
```

Example 6: Defining Multiple Indexes By Using the UNIQUE and NOMISS Options

This INDEX= data set option defines a simple index for the SSN variable and a composite index for CITY and STATE. Each variable must have a UNIQUE and NOMISS option:

```
data new(index=(ssn /unique/nomiss cityst=(city state) /unique/nomiss));
/* ... more SAS code ... */
run;
```

See Also

- [“INDEX CREATE Statement” in Base SAS Procedures Guide](#)
- [“CREATE INDEX Statement” in SAS SQL Procedure User's Guide](#)
- [“Using Indexes” in SAS Programmer's Guide: Essentials](#)

IN= Data Set Option

Creates a Boolean variable that indicates whether the data set contributed data to the current observation.

Valid in:	DATA step
Categories:	Observation Control CAS
Restriction:	Use with the SET, MERGE, MODIFY, and UPDATE statements only.

Syntax

IN=*variable*

Syntax Description

variable

names the new variable whose value indicates whether the input data set contributed data to the current observation. Within the DATA step, the value of the variable is 1 if the data set contributed to the current observation. Otherwise, the value is 0.

Details

Specify the IN= data set option in parentheses after a SAS data set name in the SET, MERGE, MODIFY, and UPDATE statements only. Values of IN= variables are available to program statements during the DATA step. These variables are not included in the SAS data set that is being created, unless they are assigned to a new variable.

When you use IN= with BY-group processing, and when a data set contributes an observation for the current BY group, the IN= value is 1. The value remains as long as that BY group is still being processed and the value is not reset by programming logic.

Example

In this example, IN= creates a new variable, OVERSEAS, that denotes international flights. The variable `I` has a value of 1 when the observation is read from the NONUSA data set. Otherwise, the variable has a value of 0. The IF-THEN statement checks the value of `I` to determine whether the data set NONUSA contributed data to the current observation. If `I=1`, the variable OVERSEAS receives an asterisk (*) as a value.

```
data allflts;
  set usa nonusa(in=i);
  by fltnum;
  if i then overseas='*';
run;
```

See Also

- [BY-group processing](#)

Statements:

- [“BY Statement” in SAS DATA Step Statements: Reference](#)
- [“MERGE Statement” in SAS DATA Step Statements: Reference](#)
- [“MODIFY Statement” in SAS DATA Step Statements: Reference](#)
- [“SET Statement” in SAS DATA Step Statements: Reference](#)
- [“UPDATE Statement” in SAS DATA Step Statements: Reference](#)

KEEP= Data Set Option

For an input data set, specifies the variables to process; for an output data set, specifies the variables to write to the data set.

Valid in: DATA step and PROC steps

Categories: Variable Control
CAS

Syntax

KEEP=*variable(s)*

Syntax Description

variable(s)

lists one or more variable names. You can list the variables in any form that SAS allows.

Details

If the KEEP= data set option is associated with an input data set, only those variables that are listed after the KEEP= data set option are available for processing. If the KEEP= data set option is associated with an output data set, all variables are available for processing. However, only the variables that are listed after the option are written to the output data set.

Comparisons

The KEEP= data set option differs from the KEEP statement in the following ways:

- In DATA steps, the KEEP= data set option can apply to input and output data sets. The KEEP statement applies only to output data sets.
- In DATA steps, when you create multiple output data sets, use the KEEP= data set option to write different variables to different data sets. The KEEP statement applies to all output data sets.
- In PROC steps, use only the KEEP= data set option, not the KEEP statement.

Example

In this example, only IDNUM and SALARY are read from PAYROLL, and they are the only variables in PAYROLL that are available for processing:

```
data bonus;  
    set payroll(keep=idnum salary);  
    bonus=salary*1.1;  
run;
```

See Also

Data Set Options:

- [“DROP= Data Set Option” on page 20](#)

Statements:

- [“KEEP Statement” in SAS DATA Step Statements: Reference](#)

LABEL= Data Set Option

Specifies a label for a SAS data set.

Valid in: DATA step and PROC steps

Categories: Data Set Control
CAS

Syntax

LABEL=*'label'*

Syntax Description

'label'

specifies a text string of up to 256 bytes. If the label text contains single quotation marks, enclose the label in double quotation marks. To remove a label from a data set, assign a blank space that is enclosed in quotation marks to the label.

You can also use two single quotation marks in the label text and enclose the string in single quotation marks.

Details

You can use the LABEL= option on input and output data sets. When you specify LABEL= on input data sets, it assigns a file label for the duration of that DATA step or PROC step. When you specify LABEL= for an output data set, the label becomes a permanent part of that file. The file can be printed using the CONTENTS or DATASETS procedure, and modified using PROC DATASETS.

A label that is assigned to a data set remains associated with that data set when you update a data set by using the APPEND procedure or the MODIFY statement. However, a label is lost if you use a data set with a previously assigned label to create a new data set in the DATA step. For example:

```
data one;  
    set one;  
run;
```

Comparisons

- The LABEL= data set option enables you to assign labels only for data sets. You can specify labels for the variables in a data set by using the LABEL statement.
- The LABEL= option enables you to assign labels to variables in the ATTRIB statement.

Example: Assigning Labels to Data Sets

```
data w2(label='1976 W2 Info, Hourly');  
/* ... more SAS code ... */  
run;  
  
data new(label='Peter's List');  
/* ... more SAS code ... */  
run;  
  
data new(label="Hillside's Daily Account");  
/* ... more SAS code ... */  
run;  
  
data sales(label='Sales For May(NE)');  
/* ... more SAS code ... */  
run;
```

See Also

Statements:

- [“ATTRIB Statement” in SAS DATA Step Statements: Reference](#)
- [“LABEL Statement” in SAS DATA Step Statements: Reference](#)
- [“MODIFY Statement” in SAS DATA Step Statements: Reference](#)

Procedures:

- [“CONTENTS Procedure” in Base SAS Procedures Guide](#)
- [“DATASETS Procedure” in Base SAS Procedures Guide](#)

MAXTABLEMEM= Data Set Option

Specifies the maximum amount of memory in bytes that each thread should allocate for in-memory blocks before converting to a memory-mapped file.

Valid in:	DATA step and PROC steps
Categories:	CAS Data Access
Default:	16M
Requirement:	You must have SAS Viya 3.3 or later.

Syntax

MAXTABLEMEM=*integer* | *integerK* | *integerM* | *integerG*

Required Arguments

integer

specifies the total number of bytes to allocate.

integerK

specifies the total number of kilobytes to allocate.

integerM

specifies the total number of megabytes to allocate.

integerG

specifies the total number of gigabytes to allocate.

Details

Integer values are always converted to the nearest whole megabyte but not less than 1 megabyte. Specifying 0 indicates to use the value from the MAXTABLEMEM= session option.

Comparisons

- The MAXTABLEMEM= data set option overrides the MAXTABLEMEM= LIBNAME statement option.

See Also

- “CASSESSOPTS= System Option” in *SAS Cloud Analytic Services: User’s Guide*
- “MAXTABLEMEM= Session Option” in *SAS Cloud Analytic Services: User’s Guide*

OBSBUF= Data Set Option

Determines the size of the view buffer for processing a DATA step view.

Valid in:	DATA step and PROC steps
Category:	Data Set Control
Restrictions:	Use in the Compute Server, but not in CAS. Valid only for a DATA step view

Syntax

OBSBUF=*n*

Syntax Description

n

specifies the number of observations that are read into the view buffer at a time.

Default 32K bytes of memory are allocated for the default view buffer, which means the default number of observations that can be read into the view buffer at one time depends on the observation length. Therefore, the default is the number of observations that can fit into 32K bytes. If the observation length is larger than 32K bytes, then only one observation can be read into the buffer at a time.

Tip To determine the observation length in bytes, use PROC CONTENTS for the DATA step view.

CAUTION **The maximum value for the OBSBUF= option depends on the amount of available memory.** If you specify a value so large that the memory allocation of the view buffer fails, an out-of-memory error results.

Details

The OBSBUF= data set option specifies the number of observations that can be read into the view buffer at a time. The *view buffer* is a segment of memory that is allocated to hold output observations that are generated from a DATA step view.

The size of the buffer determines how much data can be held in memory at one time. OBSBUF= enables you to tune the performance of reading data from a DATA step view.

The view buffer is shared between the request that opens the DATA step view (for example, a SAS procedure) and the DATA step view itself. Two computer tasks coordinate between requesting data and generating and returning the data as follows:

- When a request task (such as a PRINT procedure) requests data, task switching occurs from the request task to the view task. This action executes the DATA step view and generates the observations. The DATA step view fills the view buffer with as many observations as possible.
- When the view buffer is full, task switching occurs from the view task back to the request task in order to return the requested data. The observations are cleared from the view buffer.

The size of the view buffer determines how many generated observations can be held. The number of generated observations determines how many times the computer must switch between the request task and the view task. For example, OBSBUF=1 results in task switching for each observation. OBSBUF=10 results in 10 observations being read into the view buffer at a time. The larger the view buffer, the less task switching is needed to process a DATA step view, which can speed execution time.

To improve efficiency, determine how many observations fit in the default buffer size, and then set the view buffer so that it can hold more generated observations.

Note: Using OBSBUF= can improve processing efficiency by reducing task switching. However, the larger the view buffer size, the more time it takes to fill. This process delays the task switching from the view task back to the request task. The delay is more apparent in interactive applications. For example, when you use the Viewtable window, the larger the view buffer, the longer it takes to display the requested observations. The view buffer must be filled before even one observation is returned to the Viewtable. Before you set a very large view buffer size, consider the following information:

- the type of application that you are using to process the DATA step view
 - the amount of available memory
-

Example

For this example, the observation length is 10K. The default view buffer size, which is 32K, would result in three observations at a time to be read into the view buffer. The default view buffer size causes the execution time to be slower, because the computer must perform task switching for every three observations that are generated.

To improve performance, the OBSBUF= data set option is set to 100. This action causes 100 observations at a time to be read into the view buffer. It also reduces task switching in order to process the DATA step view with the PRINT procedure:

```
data testview / view=testview;
    ... more SAS code ...
run;
proc print data=testview (obsbuf=100);
run;
```

See Also

Data Set Options:

- [“SPILL= Data Set Option” on page 82](#)

OBS= Data Set Option

Specifies the last observation that SAS processes in a data set.

Valid in:	DATA step and PROC steps
Category:	Observation Control
Default:	MAX
Restrictions:	Use in the Compute Server, but not in CAS. Use with input data sets only. Cannot use with PROC SQL views

Syntax

OBS= *n* | *nK* | *nM* | *nG* | *nT* | *hexX* | MIN | MAX

Syntax Description

n* | *nK* | *nM* | *nG* | *nT

specifies a number to indicate when to stop processing observations, with *n* as an integer. Using one of the letter notations results in multiplying the integer by a specific value. That is, specifying K (kilo) multiplies the integer by 1,024; M (mega) multiplies by 1,048,576; G (giga) multiplies by 1,073,741,824; or T (tera) multiplies by 1,099,511,627,776. For example, a value of 20 specifies 20 observations, and a value of 3m specifies 3,145,728 observations.

Note: You can also use the KB, MB, and GB syntax notations.

hexX

specifies a number to indicate when to stop processing as a hexadecimal value. You must specify the value beginning with a number (0–9), followed by an X. For example, the hexadecimal value **F8** must be specified as **0F8x** in order to specify the decimal equivalent of 248. The value **2dx** specifies the decimal equivalent of 45.

MIN

specifies the number to indicate when to stop processing to 0. Use OBS=0 to create an empty data set that has the structure, but not the observations, of another data set.

Interaction If OBS=0 and the NOREPLACE option is in effect, SAS can still take certain actions. SAS actually executes each DATA and PROC step in the program, using no observations. For example, SAS executes procedures, such as CONTENTS and DATASETS, that process libraries or SAS data sets.

MAX

specifies the number to indicate when to stop processing to the maximum number of observations in the data set. This number can be up to the largest 8-byte, signed integer, which is $2^{63}-1$, or approximately 9.2 quintillion. This is the default.

Details

OBS= tells SAS when to stop processing observations. To determine when to stop processing, SAS uses the value for OBS= in a formula that includes the value for OBS= and the value for FIRSTOBS=.

```
(obs - firstobs) + 1 = results
```

For example, if OBS=10 and FIRSTOBS=1 (which is the default for FIRSTOBS=), the result is 10 observations. That is, $(10 - 1) + 1 = 10$. If OBS=10 and FIRSTOBS=2, the result is nine observations. That is, $(10 - 2) + 1 = 9$. OBS= is valid only when an existing SAS data set is read.

The OBS= data set option overrides the OBS= system option for the individual data set.

Comparisons

- When the OBS= data set option specifies an ending point for processing, the FIRSTOBS= data set option specifies a starting point. The two options are often used together to define a range of observations to be processed.
- The OBS= data set option enables you to select observations from SAS data sets. You can select observations to be read from external data files by using the OBS= option in the INFILE statement.

Examples

Example 1: Using OBS= to Specify When to Stop Processing Observations

This example creates a SAS data set and executes the PRINT procedure with FIRSTOBS=2 and OBS=12. The result is 11 observations. That is, $(12 - 2) + 1 = 11$. The result of OBS= appears to be the observation number that SAS processes last.

```
data Ages;
    input Name $ Age;
    datalines;
Miguel 53
Brad 27
Willie 69
Marc 50
Sylvia 40
Arun 25
Gary 40
Becky 51
Alma 39
Tom 62
Kris 66
Paul 60
Randy 43
Barbara 52
Virginia 72
;
proc print data=Ages (firstobs=2 obs=12);
run;
```

Output 2.3 PROC PRINT Output Using OBS= and FIRSTOBS=**The SAS System**

Obs	Name	Age
2	Brad	27
3	Willie	69
4	Marc	50
5	Sylvia	40
6	Arun	25
7	Gary	40
8	Becky	51
9	Alma	39
10	Tom	62
11	Kris	66
12	Paul	60

Example 2: PROC PRINT Using a WHERE Statement

This example uses the data set that was created in Example 1, which contains 15 observations.

Here is the PRINT procedure with a WHERE statement. The subset of the data results in 12 observations:

```
proc print data=Ages;
  where Age LT 65;
run;
```

Output 2.4 PROC PRINT Output Using a WHERE Statement**The SAS System**

Obs	Name	Age
1	Miguel	53
2	Brad	27
4	Marc	50
5	Sylvia	40
6	Arun	25
7	Gary	40
8	Becky	51
9	Alma	39
10	Tom	62
12	Paul	60
13	Randy	43
14	Barbara	52

Example 3: PROC PRINT Using a WHERE Statement and OBS=

Executing the PRINT procedure with the WHERE statement and OBS=10 results in 10 observations. That is, $(10 - 1) + 1 = 10$. With WHERE processing, SAS subsets the data and applies OBS= to the subset.

```
proc print data=Ages (obs=10);
    where Age LT 65;
run;
```

Output 2.5 PROC PRINT Output Using a WHERE Statement and OBS=**The SAS System**

Obs	Name	Age
1	Miguel	53
2	Brad	27
4	Marc	50
5	Sylvia	40
6	Arun	25
7	Gary	40
8	Becky	51
9	Alma	39
10	Tom	62
12	Paul	60

Example 4: PROC PRINT Using a WHERE Statement, OBS=, and FIRSTOBS=

The result of OBS= appears to be the observation number that SAS processes. If you apply FIRSTOBS=2 and OBS=10 to the subset, then the result is nine observations. That is, $(10 - 2) + 1 = 9$. OBS= is neither the observation number to end with nor how many observations to process; the value is used in the formula to determine when to stop processing.

```
proc print data=Ages (firstobs=2 obs=10);
    where Age LT 65;
run;
```

Output 2.6 PROC PRINT Output Using a WHERE Statement, OBS=, and FIRSTOBS=

The SAS System

Obs	Name	Age
2	Brad	27
4	Marc	50
5	Sylvia	40
6	Arun	25
7	Gary	40
8	Becky	51
9	Alma	39
10	Tom	62
12	Paul	60

Example 5: PROC PRINT Showing Deleted Observations

This example uses the data set that was created in Example 1, with observation 6 deleted.

Here is PROC PRINT output of the modified file:

```
proc print data=Ages;  
run;
```

Output 2.7 PROC PRINT Output Showing Observation 6 Deleted

The SAS System

Obs	Name	Age
1	Miguel	53
2	Brad	27
3	Willie	69
4	Marc	50
5	Sylvia	40
7	Gary	40
8	Becky	51
9	Alma	39
10	Tom	62
11	Kris	66
12	Paul	60
13	Randy	43
14	Barbara	52
15	Virginia	72

Example 6: PROC PRINT Using OBS=

Executing the PRINT procedure with OBS=12 results in 12 observations. That is, $(12 - 1) + 1 = 12$:

```
proc print data=Ages (obs=12);  
run;
```

Output 2.8 PROC PRINT Output Using OBS=

The SAS System

Obs	Name	Age
1	Miguel	53
2	Brad	27
3	Willie	69
4	Marc	50
5	Sylvia	40
7	Gary	40
8	Becky	51
9	Alma	39
10	Tom	62
11	Kris	66
12	Paul	60
13	Randy	43

Example 7: Using OBS= When Observations Are Deleted

The result of OBS= appears to be the observation number that SAS processes. However, if you apply FIRSTOBS=2 and OBS=12, the result is 11 observations. That is, $(12 - 2) + 1 = 11$. OBS= is neither the observation number to end with nor how many observations to process; the value is used in the formula to determine when to stop processing.

```
proc print data=Ages (firstobs=2 obs=12);
run;
```


Output 2.9 PROC PRINT Output Using OBS= and FIRSTOBS=

The SAS System

Obs	Name	Age
2	Brad	27
3	Willie	69
4	Marc	50
5	Sylvia	40
7	Gary	40
8	Becky	51
9	Alma	39
10	Tom	62
11	Kris	66
12	Paul	60
13	Randy	43

See Also

Data Set Options:

- [“FIRSTOBS= Data Set Option” on page 33](#)

Statements:

- [“INFILE Statement” in SAS DATA Step Statements: Reference](#)
- [“WHERE Statement” in SAS DATA Step Statements: Reference](#)
- [SAS DATA Step Statements: Reference](#)

System Options:

- [“OBS= System Option” in SAS System Options: Reference](#)

OUTREP= Data Set Option

Specifies the data representation for the output SAS data set.

Valid in: DATA step and PROC steps

Category:	Data Set Control
Restriction:	Use in the Compute Server, but not in CAS.
See:	OUTREP= data set option in the SAS Companion for z/OS .

Syntax

OUTREP=[SESSION](#) | *format*

Syntax Description

SESSION

specifies the data representation of the current SAS session.

The session value is also supported in the **OVERWRITE=** option in the **COPY** procedure or **COPY** statement in the **DATASETS** procedure.

format

specifies the data representation, which is the form in which data is stored in a particular operating environment. Different operating environments use different standards or conventions for storing data.

- Floating-point numbers can be represented in IEEE floating-point format or IBM floating-point format.
- Data alignment can be on a 1-byte, 4-byte, or 8-byte boundary, depending on data type requirements for the operating environment.
- Data type lengths can be 8 bits or more for a character data type, 16 bit, 32 bit, or 64 bit for an integer data type, 32 bit for a single-precision floating-point data type, and 64 bit for a double-precision floating-point data type.
- The ordering of bytes can be big Endian or little Endian.

By default, SAS creates a new SAS data set by using the data representation of the CPU that is running SAS. Specifying the **OUTREP=** option enables you to create a SAS data set with a different data representation. For example, in a UNIX environment, you can create a SAS data set that uses a Windows data representation. For more information about compatibility and data representation, see [“Cross-Environment Data Access” in SAS Programmer’s Guide: Essentials](#).

Values for **OUTREP=** are listed in the following table:

Table 2.2 Data Representation Values for **OUTREP=** Option

OUTREP= Value	Alias¹	Environment²
ALPHA_VMS_64		OpenVMS Alpha
HP_IA64	HP_ITANIUM	HP-UX for the Itanium Processor Family Architecture

OUTREP= Value	Alias ¹	Environment ²
HP_UX_32	HP_UX	HP-UX for PA-RISC
HP_UX_64		HP-UX for PA-RISC, 64-bit
LINUX_32	LINUX	Linux for Intel architecture
LINUX_X86_64		Linux for x64
LINUX_POWER_64		Linux on the Power Architecture ³
MIPS_ABI		MIPS ABI
MVS_32	MVS	31-bit SAS on z/OS
MVS_64_BFP		64-bit SAS on z/OS
RS_6000_AIX_32	RS_6000_AIX	AIX
RS_6000_AIX_64		AIX
SOLARIS_32	SOLARIS	Solaris for SPARC
SOLARIS_64		Solaris for SPARC
SOLARIS_X86_64		Solaris for x64
VMS_IA64		OpenVMS on HP Integrity
WINDOWS_32	WINDOWS	32-bit SAS on Microsoft Windows
WINDOWS_64		64-bit SAS on Microsoft Windows (for both Itanium-based systems and x64)

¹ It is recommended that you use the current values. The aliases are available for compatibility only.

² The current release of SAS does not run in many of these environments, but they are included here for completeness.

³ LINUX_POWER_64 is added in SAS Viya 3.5. It is not supported in SAS 9.

Interactions If you specify the OUTREP= option and you plan to use the file in an earlier release of SAS, you might also want to specify the EXTENDOBSCOUNTER= option. See [“EXTENDOBSCOUNTER= Data Set Option” on page 30](#).

When you use the OUTREP= data set option, the default encoding is based on the operating environment that is represented by the

OUTREP= value and the locale of the current SAS session. To assign a non-default encoding such as UTF-8, you must also specify the ENCODING= data set option. For more information about locale and encoding, see [SAS National Language Support \(NLS\): Reference Guide](#).

Tip In the COPY procedure or COPY statement of the DATASETS procedure, the default (or CLONE) behavior is to preserve many attributes, including the data representation and encoding, from the input data set. If instead you want the data representation and encoding of the SAS session that is executing the procedure, then specify OVERRIDE=(OUTREP=SESSION ENCODING=SESSION) in PROC COPY or in the COPY statement. If you do not want to preserve any attributes from the input data set, then specify NOCLONE instead.

Details

CAUTION

Transcoding could result in character data loss when encodings are incompatible. For information about encoding and transcoding, see the [SAS National Language Support \(NLS\): Reference Guide](#).

See Also

“Definitions for Cross-Environment Data Access (CEDA)” in [SAS Programmer’s Guide: Essentials](#)

POINTOBS= Data Set Option

Specifies whether SAS creates compressed data sets whose observations can be randomly accessed or sequentially accessed.

Valid in: DATA step and PROC steps

Category: Observation Control

Restrictions: Use in the Compute Server, but not in CAS.

POINTOBS= is effective only when creating a compressed data set. Otherwise, it is ignored.

Syntax

POINTOBS=YES | NO

Syntax Description

YES

causes SAS software to produce a compressed data set that might be randomly accessed by observation number. This is the default.

Here are examples of accessing data directly by observation number:

- through the POINT= option of the MODIFY and SET statements in the DATA step
- through a specific observation number with PROC FSEDIT

TIP Specifying POINTOBS=YES does not affect the efficiency of retrieving information from a data set. It does increase CPU usage by approximately 10% when creating a compressed data set and when updating or adding information to it.

NO

suppresses the ability to randomly access observations in a compressed data set by observation number.

TIP If you do not need to access data by observation number in a compressed data set, then you can improve performance by approximately 10% when you specify POINTOBS=NO in these situations:

- when you create a compressed data set
- when you update or add observations to a compressed data set

Details

REUSE=YES takes precedence over POINTOBS=YES. For example:

```
data test (compress=yes pointobs=yes reuse=yes);
```

This data set option results in a data set that has POINTOBS=NO. Because POINTOBS=YES is the default when you use compression, REUSE=YES causes POINTOBS= to change to NO.

See Also

Data Set Options:

- “COMPRESS= Data Set Option” on page 16
- “REUSE= Data Set Option” on page 77

System Options:

- “COMPRESS= System Option” in *SAS System Options: Reference*
- “REUSE= System Option” in *SAS System Options: Reference*

PW= Data Set Option

Assigns a READ, WRITE, and ALTER password to a SAS file, and enables access to a password-protected SAS file.

Valid in: DATA step and PROC steps

Category: Data Set Control

Restriction: Use in the Compute Server, but not in CAS.

Note: Check your log after this operation to ensure password security. For more information, see “[Blotting Passwords and Encryption Key Values](#)” in *SAS Programmer's Guide: Essentials*.

Syntax

PW=*password*

Syntax Description

password

must be a valid SAS name, which limits the password to eight characters and is case-insensitive.

Details

Overview of the PW= Option

The PW= option applies to all types of SAS files except catalogs. Use this option to assign a password to a SAS file or to access a password-protected SAS file.

When you replace a SAS data set that is protected by an ALTER password, the new data set inherits the ALTER password. To change the ALTER password for the new data set, use the MODIFY statement in the DATASETS procedure.

Operating Environment Information: For more information about using passwords, see the appropriate sections of the SAS documentation.

Note:

A SAS password does not control access to a SAS file beyond SAS. Use the operating system-supplied utilities and file system security controls to control access to SAS files outside of SAS.

Restriction for Metadata-Bound Libraries

You cannot use this option with data sets that are bound to metadata-bound libraries. Metadata-bound libraries are password-protected, and you must specify that password to access data sets that are bound to them. To change this password, use PROC AUTHLIB. You cannot use PROC DATASETS to add a password to a data set that is bound to a metadata-bound library. For more information, see [“AUTHLIB Procedure” in Base SAS Procedures Guide](#).

See Also

- [“Blotting Passwords and Encryption Key Values” in SAS Programmer’s Guide: Essentials](#)
- [“Manipulating Passwords” in Base SAS Procedures Guide](#)

Data Set Options:

- [“ALTER= Data Set Option” on page 9](#)
- [“ENCRYPT= Data Set Option” on page 22](#)
- [“READ= Data Set Option” on page 68](#)
- [“WRITE= Data Set Option” on page 98](#)

PWREQ= Data Set Option

Specifies whether to display a password dialog box.

Valid in:	DATA step and PROC steps
Category:	Data Set Control
Restriction:	Use in the Compute Server, but not in CAS.

Syntax

PWREQ=[YES](#) | [NO](#)

Syntax Description

YES

specifies to display a dialog box.

NO

prevents a dialog box from displaying. If a missing or invalid password is entered, the data set is not opened and an error message is written to the SAS log.

Details

In an interactive SAS session, the PWREQ= option controls whether a dialog box is displayed when an incorrect or a missing password for a password-protected SAS data set is specified. PWREQ= applies to data sets with READ=, WRITE=, or ALTER= passwords. PWREQ= is most useful in SCL applications.

See Also

Data Set Options:

- [“ALTER= Data Set Option” on page 9](#)
- [“ENCRYPT= Data Set Option” on page 22](#)
- [“PW= Data Set Option” on page 66](#)
- [“READ= Data Set Option” on page 68](#)
- [“WRITE= Data Set Option” on page 98](#)

READ= Data Set Option

Assigns a READ= password to a SAS file that prevents users from reading the file, unless they enter the password.

Valid in: DATA step and PROC steps

Category: Data Set Control

Restriction: Use in the Compute Server, but not in CAS.

Note: Check your log after this operation to ensure password security. For more information, see [“Blotting Passwords and Encryption Key Values” in SAS Programmer’s Guide: Essentials](#).

Syntax

READ=*read-password*

Syntax Description

read-password

must be a valid SAS name.

Details

Overview of the READ= Option

The READ= option applies to all types of SAS files except catalogs. Use this option to assign a password to a SAS file or to access a read-protected SAS file.

Note: A SAS password does not control access to a SAS file beyond SAS. Use the operating system-supplied utilities and file-system security controls to control access to SAS files outside of SAS.

Restriction for Metadata-Bound Libraries

You cannot use this option with data sets that are bound to metadata-bound libraries. Metadata-bound libraries are password-protected, and you must specify that password to access data sets that are bound to them. To change this password, use PROC AUTHLIB. You cannot use PROC DATASETS to add a password to a data set that is bound to a metadata-bound library. For more information, see [“AUTHLIB Procedure” in Base SAS Procedures Guide](#).

See Also

- [“Blotting Passwords and Encryption Key Values” in SAS Programmer’s Guide: Essentials](#)
- [“Manipulating Passwords” in Base SAS Procedures Guide](#)

Data Set Options:

- [“ALTER= Data Set Option” on page 9](#)
- [“ENCRYPT= Data Set Option” on page 22](#)
- [“PW= Data Set Option” on page 66](#)
- [“WRITE= Data Set Option” on page 98](#)

RENAME= Data Set Option

Changes the name of a variable.

Valid in: DATA step and PROC steps

Categories: Variable Control
CAS

Syntax

RENAME=(*old-variable-name-1=new-variable-name-1* <*old-variable-name-2=new-variable-name-2* ...>)

Syntax Description

old-name

is the variable that you want to rename.

new-name

is the new name of the variable. It must be a valid SAS name.

Details

If you use the RENAME= data set option when you create a data set, the new variable name is included in the output data set. If you use RENAME= on an input data set, the new name is used in DATA step programming statements.

If you use RENAME= on an input data set that is used in a SAS procedure, SAS changes the name of the variable in that procedure. If you use RENAME= with WHERE processing such as a WHERE statement or a WHERE= data set option, the new name is applied before the data is processed. You must use the new name in the WHERE expression.

Use RENAME= in the same DATA step with either the DROP= data set option or the KEEP= data set option. The DROP= and KEEP= data set options are applied before RENAME=. You must use the old name in the DROP= and KEEP= data set options. You cannot drop and rename the same variable in the same statement.

Note: The RENAME= data set option has an effect only on data sets that are opened in output mode.

Use the RENAME statement or the RENAME= data set option when program logic requires that you rename variables. An example is two input data sets that have

variables with the same name. To rename variables as a file management task, use the DATASETS procedure.

You must use the RENAME= data set option on the input data set or data sets to rename variables before processing begins.

Comparisons

The RENAME= data set option differs from the RENAME statement in these ways:

- You can use the RENAME= data set option, not the RENAME statement, in PROC steps.
- You must use the RENAME= data set option to rename different variables in different data sets. The RENAME statement applies to all output data sets.

Examples

Example 1: Renaming a Variable at Time of Output

This example uses RENAME= in the DATA statement to show that the variable is renamed when it is written to the output data set. The variable keeps its original name, X, during DATA step processing.

```
data one;
  input x y z;
  datalines;
24 595 439
243 343 034
;
proc print data=one;
run;

data two(rename=(x=keys));
  set one;
  z=x+y;
run;
proc print data=two;
run;
```

Output 2.10 Data Sets One and Two

The SAS System			
Obs	x	y	z
1	24	595	439
2	243	343	34

The SAS System

Obs	keys	y	z
1	24	595	619
2	243	343	586

Example 2: Renaming a Variable at Time of Input

This example renames variable X to a variable named KEYS in the SET statement, before DATA step processing.

```
data three;
    set one(rename=(x=keys));
    z=keys+y;
run;

proc print data=three;
run;
```

Output 2.11 Data Set Three

The SAS System

Obs	keys	y	z
1	24	595	619
2	243	343	586

Example 3: Renaming a Variable for a SAS Procedure with WHERE Processing

This example renames variable `score1` to a variable named `score2` for the PRINT procedure. Because the new name is applied before the data is processed, the new name must be specified in the WHERE statement.

```
data test;
    input score1;
    datalines;
26
76
86
56
;

proc print data=test (rename=(score1=score2));
    where score2 gt 75;
run;
```

Output 2.12 Data Set Test

The SAS System

Obs	score2
2	76
3	86

See Also

Data Set Options:

- [“DROP= Data Set Option” on page 20](#)
- [“KEEP= Data Set Option” on page 45](#)

Statement:

- [“RENAME Statement” in SAS DATA Step Statements: Reference](#)

Procedure:

- [“DATASETS Procedure” in Base SAS Procedures Guide](#)

REPEMPTY= Data Set Option

Specifies whether a new, empty data set can overwrite an existing SAS data set.

Valid in: DATA step and PROC steps

Category: Data Set Control

Restrictions: Use with output data sets only.
This data set option is not supported in a DATA step that runs in CAS.

Syntax

REPEMPTY=[YES](#) | [NO](#)

Syntax Description

YES

specifies that a new empty data set replaces an existing data set. This is the default.

Interaction If REPEMPTY=YES and REPLACE=NO, then the data set is not replaced.

NO

specifies that a new empty data set does not replace an existing data set.

Tips Use REPLACE=YES and REPEMPTY=NO for the convenience of replacing existing data sets with new ones that contain data.

Use REPLACE=YES and REPEMPTY=NO for the protection of not overwriting existing data sets with new empty ones that are created by mistake.

Details

Here is an example using REPEMPTY=NO and the warning that is written to the SAS log:

```
libname mylib 'c:\mydata' repempty=no;

data b;
run; /*create an empty data set*/

data mylib.a;
i=1;
run;

data mylib.a;
set b;
run; /* you will get an error that it can't replace b with an
empty data set */
```

```

25  libname mylib 'c:\mydata' repempty=no;
NOTE: Libref mylib was successfully assigned as follows:
      Engine:          V9
      Physical Name: c:\mydata
26
27  data b;run;

NOTE: The data set WORK.B has 1 observations and 0 variables.
NOTE: DATA statement used (Total process time):
      real time          0.18 seconds
      cpu time           0.01 seconds

27 !          /*create an empty data set*/
28  data mylib.a;i=1;run;

NOTE: The data set MYLIB.A has 1 observations and 1 variables.
NOTE: DATA statement used (Total process time):
      real time          0.17 seconds
      cpu time           0.01 seconds

29
30  data mylib.a;
31      set b;
32  run;

NOTE: There were 1 observations read from the data set WORK.B.
NOTE: The data set MYLIB.A has 1 observations and 0 variables.
WARNING: Data set MYLIB.A was not replaced because REPEMPTY=NO
and the replacement file is empty.
NOTE: DATA statement used (Total process time):
      real time          0.17 seconds
      cpu time           0.01 seconds

32 !          /* you will get an error that it cannot replace b with an
              empty data set */

```

For an individual data set, the REPEMPTY= data set option overrides the REPEMPTY= option in the LIBNAME statement.

Comparisons

- The REPEMPTY= and REPLACE= data set options apply to permanent and temporary SAS data sets. However, the REPLACE system option applies only to permanent SAS data sets.

See Also

Data Set Options:

- [“REPLACE= Data Set Option” on page 76](#)

Statement Options:

- [“REPEMPTY=YES | NO” in SAS Global Statements: Reference](#)

System Options:

- [“REPLACE System Option” in SAS System Options: Reference](#)

REPLACE= Data Set Option

Specifies whether a new SAS data set that contains data can overwrite an existing data set.

Valid in: DATA step and PROC steps

Categories: Data Set Control
CAS

Restrictions: Use with output data sets only.
This option is valid only when creating a SAS data set.

Note: When you use the OUT2= *PermanentLibrary_ALL_* option within PROC CONTENTS or PROC DATASETS with the CONTENTS statement, you must also set the REPLACE=YES data set option or the REPLACE system option.

Syntax

REPLACE=[NO](#) | [YES](#)

Syntax Description

NO

specifies that a new data set does not replace an existing data set.

YES

specifies that a new data set replaces an existing data set.

Details

- The REPLACE= data set option overrides the REPLACE system option for the individual data set.
- The REPLACE system option applies only to permanent SAS data sets.

Use REPLACE=YES and REPEMPTY=NO for the convenience of replacing existing data sets with new ones that contain data.

Example

Using the REPLACE= data set option in this DATA statement prevents SAS from replacing a SAS data set named ONE in a library referenced by MYLIB:


```
data mylib.one(replace=no);
```

SAS writes a message to the log that the file has not been replaced.

See Also

System Options:

- [“REPLACE System Option” in SAS System Options: Reference](#)

REUSE= Data Set Option

Specifies whether new observations can be written to available space in compressed SAS data sets.

Valid in:	DATA step and PROC steps
Category:	Data Set Control
Restrictions:	Use in the Compute Server, but not in CAS. Use with output data sets only.

Syntax

REUSE=[NO](#) | [YES](#)

Syntax Description

NO

does not track and reuse space in compressed data sets. New observations are appended to the existing data set. Specifying the NO argument results in less efficient data storage if you delete or update many observations in the SAS data set.

You can use procedures such as APPEND and FSEDIT that add observations to the end of SAS data sets with compressed data sets.

YES

tracks and reuses space in compressed SAS data sets. New observations are inserted in the space that is available when other observations are updated or deleted.

REUSE=YES causes new observations to be added wherever there is space in the file, not necessarily at the end of the file.

Details

By default, new observations are appended to existing compressed data sets. To track and reuse available space by deleting or updating other observations, use the REUSE= data set option when you create a compressed SAS data set.

Use REUSE= only when you are creating new data sets with the COMPRESS=YES data set option or system option.

The REUSE= data set option overrides the REUSE= system option.

REUSE=YES takes precedence over POINTOBS=YES. For example, the following statement results in a data set that has POINTOBS=NO:

```
data test (compress=yes pointobs=yes reuse=yes);
```

Because POINTOBS=YES is the default when you use compression, REUSE=YES causes POINTOBS= to change to NO.

See Also

Data Set Options:

- [“COMPRESS= Data Set Option” on page 16](#)

System Options:

- [“REUSE= System Option” in SAS System Options: Reference](#)

ROLE= Data Set Option

Identifies the fact table for a star schema join.

Valid in: PROC SQL

Category: Data Set Control

Restrictions: Use in the Compute Server, but not in CAS.
Use with input data sets only.

Syntax

ROLE=[FACT](#) | [DIMENSION](#) | [DIM](#)

Syntax Description

FACT

identifies the SAS data set as the fact table for a star schema.

DIMENSION | DIM

identifies the SAS data set as a dimension table for a star schema.

Details

A *star schema* is an arrangement of several tables in which a large fact table is joined to several dimension tables. For example, you can join SAS data sets by using SQL procedure syntax to create a star schema.

To improve the performance of the application that processes the joined tables, specify the `ROLE=` data set option. For example, specify `ROLE=FACT` to designate the specific fact table. You can also specify `ROLE=DIMENSION` to designate each dimension table.

Because the role a table plays can change between queries, the `ROLE=` specification is in effect for the current step only and is not stored with the data set.

Example: Designating the Fact Table

In this example, the `ROLE=` data set option improves the performance of PROC SQL. `ORDERS` is the fact table, and `PRODUCT`, `PERIOD`, and `CUSTOMER` are dimension tables.

```
proc sql;
  select orders.Order_Total
  from orders (role=fact), product, period, customer
  where orders.Product_ID = product.Product_ID
        and orders.Period_ID = period.Period_ID
        and orders.Customer_ID = customer.Customer_ID
        and product.Product_Name = "camera"
        and period.Period_Name = "1997"
        and customer.Customer_Name = "Walmart";
quit;
```

See Also

[“SQL Procedure” in SAS SQL Procedure User's Guide](#)

SORTEDBY= Data Set Option

Specifies how a data set is currently sorted.

Valid in: DATA step and PROC steps

Category: Data Set Control

Restriction: Use in the Compute Server, but not in CAS.

Syntax

SORTEDBY=*by-clause***</ collate-name> | _NULL_**

Syntax Description

***by-clause* < / collate-name>**

indicates how the data is currently sorted.

by-clause

names the variables and options that you use in a BY statement in a PROC SORT step.

collate-name

names the collating sequence that is used for the sort. By default, the collating sequence is that of your operating environment. A slash (/) must precede the collating sequence.

Operating Environment Information: For more information about collating sequences, see the SAS documentation for your operating environment.

NULL

removes any existing sort indicator.

Details

SAS determines whether a data set is already sorted by the key variable or variables in ascending order by checking the sort indicator. The sort indicator is stored in the data set descriptor information and is set from a previous sort. For more information about how the sort indicator is used and how it improves performance, see [“Example: View Sort Information for a Data Set” in SAS Programmer’s Guide: Essentials](#) and [“SORTVALIDATE System Option” in SAS System Options: Reference](#).

This example of the CONTENTS procedure indicates that the data set was sorted using the SORTEDBY= data set option.

```
Sort Information
Sortedby var1
Validated NO
Character Set ANSI
```

Comparisons

- The CONTENTS statement in the DATASETS procedure indicates how a data set is sorted.

- The SORTEDBY= option indicates how the data is sorted, but does not cause a data set to be sorted.

Example

This example uses the SORTEDBY= data set option to specify how the data is currently sorted. The data set ORDERS is sorted by PRIORITY and by the descending values of INDATE. Once the data set is created, the sort indicator is stored with it. These statements create the data set ORDERS and record the sort indicator:

```
libname mylib 'c:\mydata';

options yearcutoff=1926;
data mylib.orders(sortedby=priority
                  descending indate);
  input priority 1. +1 indate date7.
        +1 office $ code $;
  format indate date7.;
  datalines;
1 03may01 CH J8U
1 21mar01 LA M91
1 01dec00 FW L6R
1 27feb99 FW Q2A
2 15jan08 FW I9U
2 09jul99 CH P3Q
3 08apr10 CH H5T
3 31jan12 FW D2W
;
proc contents data=mylib.orders;
run;
quit;
```

Output 2.13 PROC CONTENTS Sort Information

Sort Information	
Sortedby	priority DESCENDING indate
Validated	NO
Character Set	ANSI

See Also

- [“DATASETS Procedure” in Base SAS Procedures Guide](#)
- [“SORT Procedure” in Base SAS Procedures Guide](#)
- [“SQL Procedure” in SAS SQL Procedure User’s Guide](#)

SPILL= Data Set Option

Specifies whether to create a spill file for non-sequential processing of a DATA step view.

Valid in:	DATA step and PROC steps
Category:	Data Set Control
Restrictions:	Use in the Compute Server, but not in CAS. Valid only for a DATA step view

Syntax

SPILL=YES | NO

Syntax Description

YES

creates a spill file for non-sequential processing of a DATA step view. This is the default.

Interaction A spill file is never created for sequential processing of a DATA step view.

NO

does not create a spill file or reduce the size of a spill file.

Interaction For direct (random) access, a spill file is always created even if you specify SPILL=NO.

Note If you do not have enough disk space to accommodate a resulting spill file from a DATA step view that generates a large amount of data, specify SPILL=NO.

Tip For SAS procedures that process BY-group data, consider specifying SPILL=NO in order to write only the current BY group to the spill file.

Details

When a DATA step view is opened for non-sequential processing, a spill file is created by default. The *spill file* contains the observations that are generated by a DATA step view. Subsequent requests for data read the observations from the spill file rather than execute the DATA step view again. The spill file is a temporary file in the Work library.

Non-sequential processing includes the following access methods, which are supported by several SAS statements and procedures. How the SPILL= data set option operates with each of the access methods is described here:

random access

retrieves observations directly either by an observation number or by the value of one or more variables through an index without reading all observations sequentially. Whether SPILL=YES or SPILL=NO, a spill file is always created, because the processing time to restart a DATA step view for each observation is significant.

BY-group access

uses a BY statement to process observations that are ordered, grouped, or indexed according to the values of one or more variables. SPILL=YES creates a spill file the size of all the data that is requested from the DATA step view. SPILL=NO writes only the current BY group to the spill file. The size of a spill file depends on the size of a BY group.

two-pass access

performs multiple sequential passes through the data. With SPILL=NO, no spill file is created. Instead, after the first pass through the data, the DATA step view is restarted for each subsequent pass through the data. If small amounts of data are returned by the DATA step view for each restart, the processing time to restart the view might be significant.

Note: With SPILL=NO, subsequent passes through the data could result in generating different data. Some processing might require using a spill file. For example, results from using random functions and computing values that are based on the current time of day could affect the data.

Examples

Example 1: Using a Spill File for a Small Number of Large BY Groups

This example creates a DATA step view that generates a large amount of random data and uses the UNIVARIATE procedure with a BY statement. The example illustrates the effects of SPILL= with a small number of large BY groups.

With SPILL=YES, all observations that are requested from the DATA step view are written to the spill file. With SPILL=NO, only the observations that are in the current BY group are written to the spill file. The output messages that are produced by this example show that the size of the spill file is reduced with SPILL=NO. However, the time it takes to truncate the spill file for each BY group might add to the overall processing time for the DATA step view.

```
options msglevel=i;
data vw_few_large / view=vw_few_large;
  drop i;
  do byval = 'Group A', 'Group B', 'Group C';
    do i = 1 to 500000;
      r = ranuni(4);
```

```
        output;  
    end;  
end;  
run;  
proc univariate data=vw_few_large (spill=yes) noprint;  
    var r;  
    by byval;  
run;  
proc univariate data=vw_few_large (spill=no) noprint;  
    var r;  
    by byval;  
run;
```


Example Code 2.1 SAS Log Output

```

1  options msglevel=i;
2  data vw_few_large / view=vw_few_large;
3      drop i;
4
5      do byval = 'Group A', 'Group B', 'Group C';
6          do i = 1 to 500000;
7              r = ranuni(4);
8              output;
9          end;
10     end;
11 run;
NOTE: DATA STEP view saved on file WORK.VW_FEW_LARGE.
NOTE: A stored DATA STEP view cannot run under a different operating system.
NOTE: DATA statement used (Total process time):
      real time           21.57 seconds
      cpu time            1.31 seconds
12  proc univariate data=vw_few_large (spill=yes) noprint;
INFO: View WORK.VW_FEW_LARGE open mode: BY-group rewind.
13      var r;
14      by byval;
15  run;
INFO: View WORK.VW_FEW_LARGE opening spill file for output observations.
INFO: View WORK.VW_FEW_LARGE deleting spill file. File size was 22506120 bytes.
NOTE: View WORK.VW_FEW_LARGE.VIEW used (Total process time):
      real time           40.68 seconds
      cpu time            12.71 seconds
NOTE: PROCEDURE UNIVARIATE used (Total process time):
      real time           57.63 seconds
      cpu time            13.12 seconds
16
17  proc univariate data=vw_few_large (spill=no) noprint;
INFO: View WORK.VW_FEW_LARGE open mode: BY-group rewind.
18      var r;
19      by byval;
20  run;
INFO: View WORK.VW_FEW_LARGE opening spill file for output observations.
INFO: View WORK.VW_FEW_LARGE truncating spill file. File size was 7502040 bytes.
NOTE: The above message was for the following by-group:
      byval=Group A
INFO: View WORK.VW_FEW_LARGE truncating spill file. File size was 7534800 bytes.
NOTE: The above message was for the following by-group:
      byval=Group B
INFO: View WORK.VW_FEW_LARGE truncating spill file. File size was 7534800 bytes.
NOTE: The above message was for the following by-group:
      byval=Group C
INFO: View WORK.VW_FEW_LARGE deleting spill file. File size was 32760 bytes.
NOTE: View WORK.VW_FEW_LARGE.VIEW used (Total process time):
      real time           11.03 seconds
      cpu time            10.95 seconds
NOTE: PROCEDURE UNIVARIATE used (Total process time):
      real time           11.04 seconds
      cpu time            10.96 seconds

```

Example 2: Using a Spill File for a Large Number of Small BY Groups

This example creates a DATA step view that generates a large amount of random data and uses the UNIVARIATE procedure with a BY statement. This example illustrates the effects of SPILL= with a large number of small BY groups.

With SPILL=YES, all observations that are requested from the DATA step view are written to the spill file. With SPILL=NO, only the observations that are in the

current BY group are written to the spill file. The output messages that are produced by this example show that the size of the spill file is reduced with SPILL=NO. Small BY groups result in large space savings.

```
options msglevel=i;
data vw_many_small / view=vw_many_small;
  drop i;
  do byval = 1 to 100000;
    do i = 1 to 5;
      r = ranuni(4);
      output;
    end;
  end;
run;
proc univariate data=vw_many_small (spill=yes) noprint;
  var r;
  by byval;
run;
proc univariate data=vw_many_small (spill=no) noprint;
  var r;
  by byval;
run;
```

Example Code 2.2 SAS Log Output

```

1  options msglevel=i;
2  data vw_many_small / view=vw_many_small;
3      drop i;
4
5      do byval = 1 to 100000;
6          do i = 1 to 5;
7              r = ranuni(4);
8              output;
9          end;
10     end;
11 run;
NOTE: DATA STEP view saved on file WORK.VW_MANY_SMALL.
NOTE: A stored DATA STEP view cannot run under a different operating system.
NOTE: DATA statement used (Total process time):
      real time          0.56 seconds
      cpu time           0.03 seconds
12  proc univariate data=vw_many_small (spill=yes) noprint;
INFO: View WORK.VW_MANY_SMALL open mode: BY-group rewind.
13      var r;
14      by byval;
15  run;
INFO: View WORK.VW_MANY_SMALL opening spill file for output observations.
INFO: View WORK.VW_MANY_SMALL deleting spill file. File size was 8024240 bytes.
NOTE: View WORK.VW_MANY_SMALL.VIEW used (Total process time):
      real time          30.73 seconds
      cpu time           29.59 seconds
NOTE: PROCEDURE UNIVARIATE used (Total process time):
      real time          30.96 seconds
      cpu time           29.68 seconds
16
17  proc univariate data=vw_many_small (spill=no) noprint;
INFO: View WORK.VW_MANY_SMALL open mode: BY-group rewind.
18      var r;
19      by byval;
20  run;
INFO: View WORK.VW_MANY_SMALL opening spill file for output observations.
INFO: View WORK.VW_MANY_SMALL truncating spill file. File size was 65504 bytes.
NOTE: The above message was for the following by-group:
      byval=410
INFO: View WORK.VW_MANY_SMALL truncating spill file. File size was 65504 bytes.
NOTE: The above message was for the following by-group:
      byval=819
INFO: View WORK.VW_MANY_SMALL truncating spill file. File size was 65504 bytes.
NOTE: The above message was for the following by-group:
      byval=1229

.
. Deleted many INFO and NOTE messages for BY groups
.
INFO: View WORK.VW_MANY_SMALL truncating spill file. File size was 65504 bytes.
NOTE: The above message was for the following by-group:
      byval=99894
INFO: View WORK.VW_MANY_SMALL deleting spill file. File size was 32752 bytes.
NOTE: View WORK.VW_MANY_SMALL.VIEW used (Total process time):
      real time          29.43 seconds
      cpu time           28.81 seconds
NOTE: PROCEDURE UNIVARIATE used (Total process time):
      real time          29.43 seconds
      cpu time           28.81 seconds

```

Example 3: Using a Spill File with Two-Pass Access

This example creates a DATA step view that generates a large amount of random data and uses the TRANSPOSE procedure. The example illustrates the effects of SPILL= with a procedure that requires two-pass access processing.

When PROC TRANSPOSE processes a DATA step view, the procedure must make two passes through the observations that the view generates. The first pass counts the number of observations and the second pass performs the transposition. With SPILL=YES, a spill file is created during the first pass, and the second pass reads the observations from the spill file. With SPILL=NO, a spill file is not created. After the first pass, the DATA step view is restarted.

The first TRANSPOSE procedure does not include the SPILL= data set option, even though a spill file is used by default. A SAS log message about the Open mode is not displayed.

```
options msglevel=i;
data vw_transpose/view=vw_transpose;
  drop i j;
  array x[10000];
  do i = 1 to 10;
    do j = 1 to dim(x);
      x[j] = ranuni(4);
    end;
    output;
  end;
run;
proc transpose data=vw_transpose out=transposed;
run;
proc transpose data=vw_transpose(spill=yes) out=transposed;
run;
proc transpose data=vw_transpose(spill=no) out=transposed;
run;
```

Example Code 2.3 SAS Log Output

```

1  options msglevel=i;
2  data vw_transpose/view=vw_transpose;
3      drop i j;
4      array x[10000];
5      do i = 1 to 10;
6          do j = 1 to dim(x);
7              x[j] = ranuni(4);
8          end;
9          output;
10     end;
11 run;
NOTE: DATA STEP view saved on file WORK.VW_TRANSPOSE.
NOTE: A stored DATA STEP view cannot run under a different operating system.
NOTE: DATA statement used (Total process time):
      real time          0.68 seconds
      cpu time           0.18 seconds
12  proc transpose data=vw_transpose out=transposed;
13  run;
INFO: View WORK.VW_TRANSPOSE opening spill file for output observations.
INFO: View WORK.VW_TRANSPOSE deleting spill file. File size was 880000 bytes.
NOTE: View WORK.VW_TRANSPOSE.VIEW used (Total process time):
      real time          2.37 seconds
      cpu time           1.17 seconds
NOTE: There were 10 observations read from the data set WORK.VW_TRANSPOSE.
NOTE: The data set WORK.TRANSPOSED has 10000 observations and 11 variables.
NOTE: PROCEDURE TRANSPOSE used (Total process time):
      real time          4.17 seconds
      cpu time           1.51 seconds
14  proc transpose data=vw_transpose (spill=yes) out=transposed;
INFO: View WORK.VW_TRANSPOSE open mode: sequential.
15  run;
INFO: View WORK.VW_TRANSPOSE reopen mode: two-pass.
INFO: View WORK.VW_TRANSPOSE opening spill file for output observations.
INFO: View WORK.VW_TRANSPOSE deleting spill file. File size was 880000 bytes.
NOTE: View WORK.VW_TRANSPOSE.VIEW used (Total process time):
      real time          0.95 seconds
      cpu time           0.92 seconds
NOTE: There were 10 observations read from the data set WORK.VW_TRANSPOSE.
NOTE: The data set WORK.TRANSPOSED has 10000 observations and 11 variables.
NOTE: PROCEDURE TRANSPOSE used (Total process time):
      real time          1.01 seconds
      cpu time           0.98 seconds
16  proc transpose data=vw_transpose (spill=no) out=transposed;
INFO: View WORK.VW_TRANSPOSE open mode: sequential.
17  run;
INFO: View WORK.VW_TRANSPOSE reopen mode: two-pass.
INFO: View WORK.VW_TRANSPOSE restarting for another pass through the data.
NOTE: View WORK.VW_TRANSPOSE.VIEW used (Total process time):
      real time          1.34 seconds
      cpu time           1.32 seconds
NOTE: The View WORK.VW_TRANSPOSE was restarted 1 times. The following view statistics
      only apply to the last view restart.
NOTE: There were 10 observations read from the data set WORK.VW_TRANSPOSE.
NOTE: The data set WORK.TRANSPOSED has 10000 observations and 11 variables.
NOTE: PROCEDURE TRANSPOSE used (Total process time):
      real time          1.42 seconds
      cpu time           1.40 seconds

```

See Also

Data Set Options:

- [“OBSBUF= Data Set Option” on page 50](#)

TOBSNO= Data Set Option

Specifies the number of observations to send in a client/server transfer.

Valid in: DATA step and PROC steps

Category: Data Set Control

Restrictions: Use in the Compute Server, but not in CAS.

The TOBSNO= option is valid only for data sets that are accessed through a SAS server by using the REMOTE engine.

Syntax

TOBSNO=*n*

Syntax Description

n

specifies the number of observations to be transmitted.

Details

If the TOBSNO= option is not specified, its value is calculated based on the observation length and the size of the server's transmission buffers. This action is specified by the PROC SERVER statement TBUFSIZE= option.

The TOBSNO= option is valid only for data sets that are accessed through a SAS server via the REMOTE engine. If this option is specified for a data set that is opened for update or accessed via another engine, it is ignored.

See Also

[“FOPEN Function” in SAS Functions and CALL Routines: Reference](#)

TYPE= Data Set Option

Specifies the data set type for a specially structured SAS data set.

Valid in:	DATA step and PROC steps
Category:	Data Set Control
Restriction:	Use in the Compute Server, but not in CAS.

Syntax

TYPE=*data-set-type*

Syntax Description

data-set-type

specifies the special type of data set.

Details

Use the TYPE= data set option in a DATA step to perform these tasks:

- to create a special SAS data set in the proper format
- to identify the special type of SAS data set in a procedure statement

You can use the CONTENTS procedure to determine the type of data set.

Most SAS data sets do not have a specified type. However, there are several specially structured SAS data sets that are used by some SAS/STAT procedures. These SAS data sets contain special variables and observations, and they are usually created by SAS statistical procedures. Because most of the special SAS data sets are used with SAS/STAT software, they are described in the *SAS/STAT User's Guide*. Some of the special data sets are CORR, COV, SSPC, EST, and FACTOR.

Additional values are available in other SAS software products and are described in the appropriate documentation.

Note: If you use a DATA step with a SET statement to modify a special SAS data set, you must specify the TYPE= option in the DATA statement. The *data-set-type* is not automatically copied to the data set that is created.

See Also

“CONTENTS Procedure” in *Base SAS Procedures Guide*

USEDIRECTIO= Data Set Option: Linux

Turns on direct file I/O for a library that contains the file to which the ENABLEDIRECTIO option has been applied.

Valid in:	DATA step
Category:	Data Set Control
Default:	Off
Engine:	V9, V8
Operating environment:	To use this option, you must also use the ENABLEDIRECTIO option in the LIBNAME statement where the libref was assigned.

Syntax

USEDIRECTIO=YES | NO

Required Argument

YES | NO

specifies whether to turn on the USEDIRECTIO= option.

Details

The USEDIRECTIO= data set option turns on direct file I/O for a data set that is listed in a DATA statement. The associated libref must have been defined with the ENABLEDIRECTIO option in the LIBNAME statement.

Using ENABLEDIRECTIO in a LIBNAME statement makes direct file I/O possible for data sets in that library. Direct I/O itself is not turned on. You must use the USEDIRECTIO= option to produce direct file I/O.

You can turn on direct file I/O in two ways:

- Use both the ENABLEDIRECTIO and USEDIRECTIO= options in the LIBNAME statement:

```
libname libref-name '.' ENABLEDIRECTIO USEDIRECTIO=yes;
```

In this case, SAS uses direct file I/O on all SAS I/O data sets that are opened using the libref *libref-name*.

- Use ENABLEDIRECTIO in the LIBNAME statement and use USEDIRECTIO= in a DATA statement:

```
libname libref-name '.' ENABLEDIRECTIO;
data libref-name.data-set-name (USEDIRECTIO=yes);
```

In this case, *libref-name.data-set-name* is opened for direct file I/O. Other SAS I/O data sets referenced by *libref-name* will not use direct file I/O.

USEDIRECTIO= by itself has no effect. Neither of the following statements open a data set for direct file I/O:

```
libname libref-name '.' USEDIRECTIO=yes;
data libref-name.data-set-name (USEDIRECTIO=yes);
```

Example

The following example uses the ENABLEDIRECTIO LIBNAME option to enable files that are associated with the libref *test* to be opened for direct I/O. The USEDIRECTIO= data set option opens *test.file1* for direct I/O. *test.file2* is not opened for direct I/O.

```
LIBNAME test '.' ENABLEDIRECTIO;

data test.file1(USEDIRECTIO=yes);
/* ... more SAS code ... */
run;
data test.file2;
/* ... more SAS code ... */
run;
```

WHERE= Data Set Option

Specifies specific conditions to use to select observations from a SAS data set.

Valid in:	DATA step and PROC steps
Categories:	Observation Control CAS
Restriction:	Cannot be used with the POINT= option in the SET and MODIFY statements.

Syntax

WHERE=(*where-expression-1* <logical-operator *where-expression-2*>)

Syntax Description

where-expression

is an arithmetic or logical expression that consists of a sequence of operators, operands, and SAS functions. An *operand* is a variable, a SAS function, or a constant. An *operator* is a symbol that requests a comparison, logical operation, or arithmetic calculation. The expression must be enclosed in parentheses.

logical-operator

can be AND, AND NOT, OR, or OR NOT.

Details

The WHERE= data set option is supported in the SET statement, which is on the input data set. The data set option is in a DATA step that runs in CAS:

```
/* Valid - DATA step runs in CAS */
libname mycas cas;

data mycas.cars;
    set mycas.cars(where=(weight>6000));
run;
```

The WHERE= data set option is supported in the DATA statement, which is on the output CAS table. The data set option is in a DATA step that runs in SAS. For example, if you are using the DATA step to load a SAS data set to a CAS table and you specify the WHERE= option on the output CAS table, then the DATA step runs in SAS and the WHERE= option is supported:

```
/* Valid - DATA step runs in SAS */
libname mycas cas;

data mycas.cars(where=(weight>6000));
    set sashelp.cars;
run;
```

Use the WHERE= data set option with an input data set to select observations that meet the condition that is specified in the WHERE expression. SAS brings the observations into the DATA or PROC step for processing. Selecting observations that meet the conditions of the WHERE expression is the first operation SAS performs in each iteration of the DATA step.

You can also select observations that are written to an output data set. In general, selecting observations at the point of input is more efficient than selecting them at the point of output. However, there are some cases when selecting observations at the point of input is not practical or not possible.

You can apply OBS= and FIRSTOBS= processing to WHERE processing. These data sets options are not valid in a DATA step running on the CAS server.

You cannot use the WHERE= data set option with the POINT= option in the SET and MODIFY statements. The POINT= data sets option is not valid in a DATA step running on the CAS server.

You can use both the WHERE= data set option and the WHERE statement in the same DATA step. SAS ignores the WHERE statement for data sets with the WHERE= data set option. However, you can use the WHERE= data set option with the WHERE command in SAS/FSP software.

Note: Using indexed SAS data sets can improve performance significantly when you are using WHERE expressions to access a subset of the observations in a SAS data set. Indexes are not supported on the CAS server.

Comparisons

- The WHERE statement applies to all input data sets, whereas the WHERE= data set option selects observations only from the data set for which it is specified.
- The DROP= and KEEP= data set options select variables for processing, whereas the WHERE= data set option selects observations.

Examples

Example 1: Selecting Observations from an Input Data Set

This example uses the WHERE= data set option to subset the SALES data set as it is read into another data set:

```
data whizmo;
    set sales(where=(product='whizmo'));
run;
```

Output 2.14 Input Data Set Observations

Whizmo Data Set			
Obs	product	sales	store
1	whizmo	234	mountain
2	whizmo	273	lakeside
3	whizmo	234	parkview
4	whizmo	233	central

Example 2: Selecting Observations from an Output Data Set

This example uses the WHERE= data set option to subset the SALES output data set:

```
data whizmo(where=(product='whizmo'));
  set sales;
run;
```

Output 2.15 Output Data Set Observations

Whizmo Data Set			
Obs	product	sales	store
1	whizmo	234	mountain
2	whizmo	273	lakeside
3	whizmo	234	parkview
4	whizmo	233	central

See Also

- [“WHERE Statement” in SAS DATA Step Statements: Reference](#)
- [WHERE-Expression Processing](#)

WHEREUP= Data Set Option

Specifies whether to evaluate new observations and modified observations against a WHERE expression.

Valid in: DATA step and PROC steps

Category: Observation Control

Restriction: Use in the Compute Server, but not in CAS.

Syntax

WHEREUP=[NO](#) | [YES](#)

Syntax Description

NO

does not evaluate added observations and modified observations against a WHERE expression.

YES

evaluates added observations and modified observations against a WHERE expression.

Details

Specify WHEREUP=YES when you want any added observations or modified observations to match a specified WHERE expression.

Examples

Example 1: Accepting Updates That Do Not Match the WHERE Expression

This example shows how WHEREUP= allows observations to be updated and added even though the modified observation does not match the WHERE expression:

```
data a;
  x=1;
  output;
  x=2;
  output;
run;
data a;
  modify a(where=(x=1) whereup=no);
  x=3;
  replace; /* Update does not match WHERE expression */
  output; /* Add does not match WHERE expression */
run;
```

In this example, SAS updates the observation and adds the new observation to the data set.

Example 2: Rejecting Updates That Do Not Match the WHERE Expression

In this example, WHEREUP= does not allow observations to be updated or added when the update and the addition do not match the WHERE expression:

```
data a;
  x=1;
  output;
  x=2;
  output;
run;
data a;
  modify a(where=(x=1) whereup=yes);
  x=3;
  replace; /* Update does not match WHERE expression */
```

```
output; /* Add does not match WHERE expression */
run;
```

In this example, SAS does not update the observation nor does it add the new observation to the data set.

See Also

Data Set Option:

- [“WHERE= Data Set Option” on page 93](#)

WRITE= Data Set Option

Assigns a WRITE= password to a SAS file that prevents users from writing to a file, unless the users enter the password.

Valid in:	DATA step and PROC steps
Category:	Data Set Control
Restriction:	Use in the Compute Server, but not in CAS.
Note:	Check your log after this operation to ensure password security. For more information, see “Blotting Passwords and Encryption Key Values” in SAS Programmer’s Guide: Essentials .

Syntax

WRITE=*write-password*

Syntax Description

write-password
must be a valid SAS name.

Details

Overview of the WRITE= Option

The WRITE= option applies to all types of SAS files except catalogs. You can use this option to assign a password to a SAS file or to access a write-protected SAS file.

Note: A SAS password does not control access to a SAS file beyond SAS. You should use the operating system-supplied utilities and file-system security controls in order to control access to SAS files outside of SAS.

Restriction for Metadata-Bound Libraries

You cannot use this option with data sets that are bound to metadata-bound libraries. Metadata-bound libraries are password-protected, and you must specify that password to access data sets that are bound to them. To change this password, use PROC AUTHLIB. You cannot use PROC DATASETS to add a password to a data set that is bound to a metadata-bound library. For more information, see [“AUTHLIB Procedure” in *Base SAS Procedures Guide*](#).

See Also

- [“Blotting Passwords and Encryption Key Values” in *SAS Programmer’s Guide: Essentials*](#)
- [“Manipulating Passwords” in *Base SAS Procedures Guide*](#)

Data Set Options:

- [“ALTER= Data Set Option” on page 9](#)
- [“ENCRYPT= Data Set Option” on page 22](#)
- [“PW= Data Set Option” on page 66](#)
- [“READ= Data Set Option” on page 68](#)

